

Compositional Conditioning Consistency Model

賴柏宏 Oral Defense
College of Artificial Intelligence,
National Yang Ming Chiao Tung University
2025 07.02

Outline

1

Introduction
Page 3-7

2

Related Works and Preliminaries
Page 8-21

3

Proposed Methods
Page 22-33

4

Experiments
Page 34-43

5

Conclusion
Page 45

6

References and Appendix
Page 46-53

1

Introduction

Our Contributions

1. We proposed **CCCM** -- the first consistency model capable of compositional zero-shot generation, effectively transferring CCDM's unseen image generation ability into 2–4 step
2. **Modified consistency distillation**, combining teacher-predicted and forward-process-formulated supervision. Three fusion strategies: Switch, Step Fuse, and Loss Fuse
3. CCCM achieves superior FID scores and maintains zero-shot accuracy despite requiring only a fraction of CCDM's sampling steps.

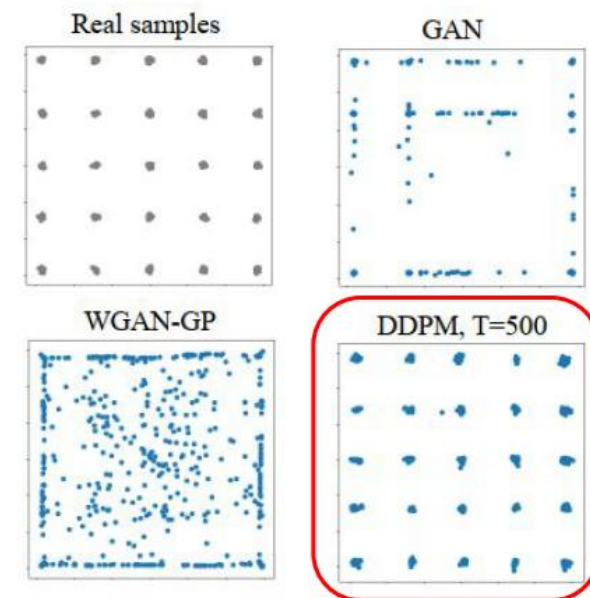
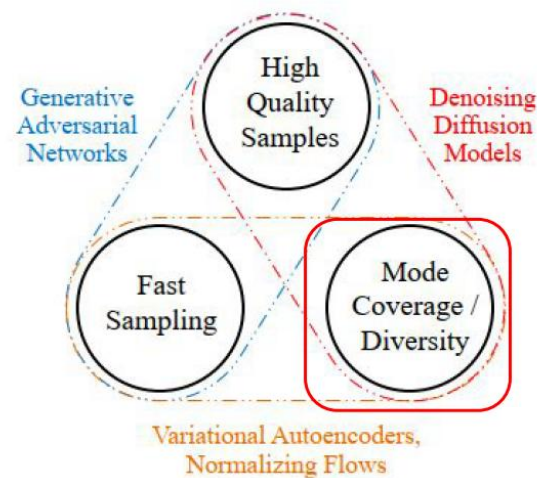
Backgrounds

- Are there neural networks capable of generating unseen classes?



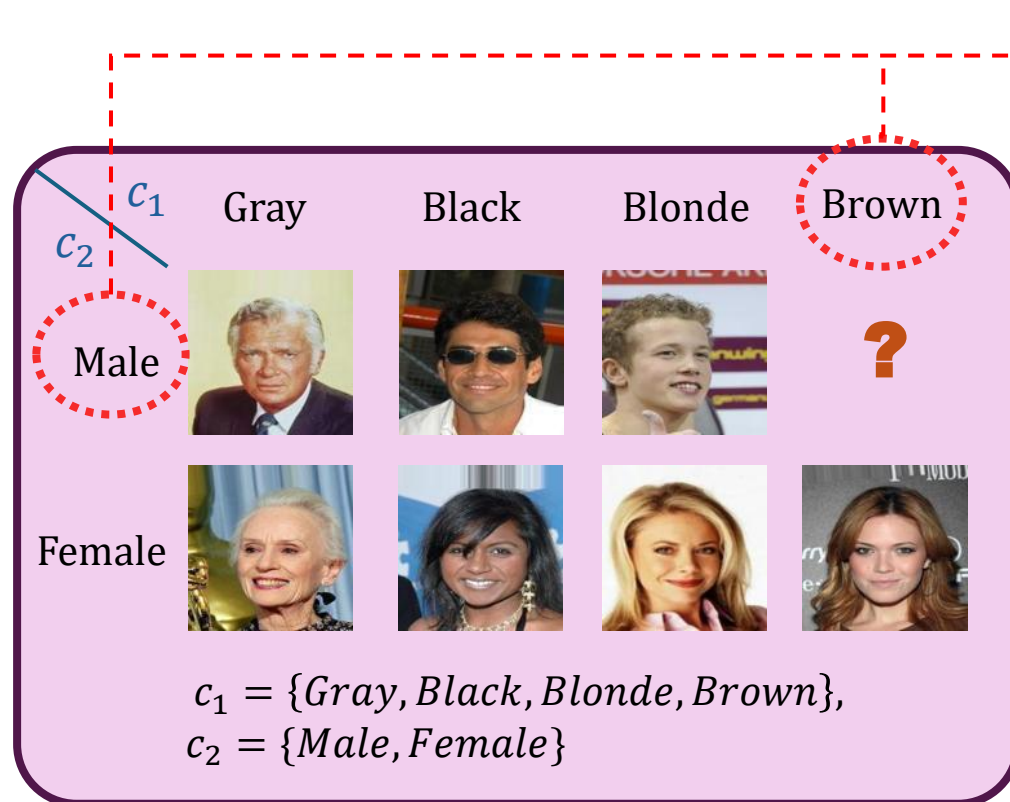
Dataset divided into 8 classes with compositional labels

Which generative model should we use ?



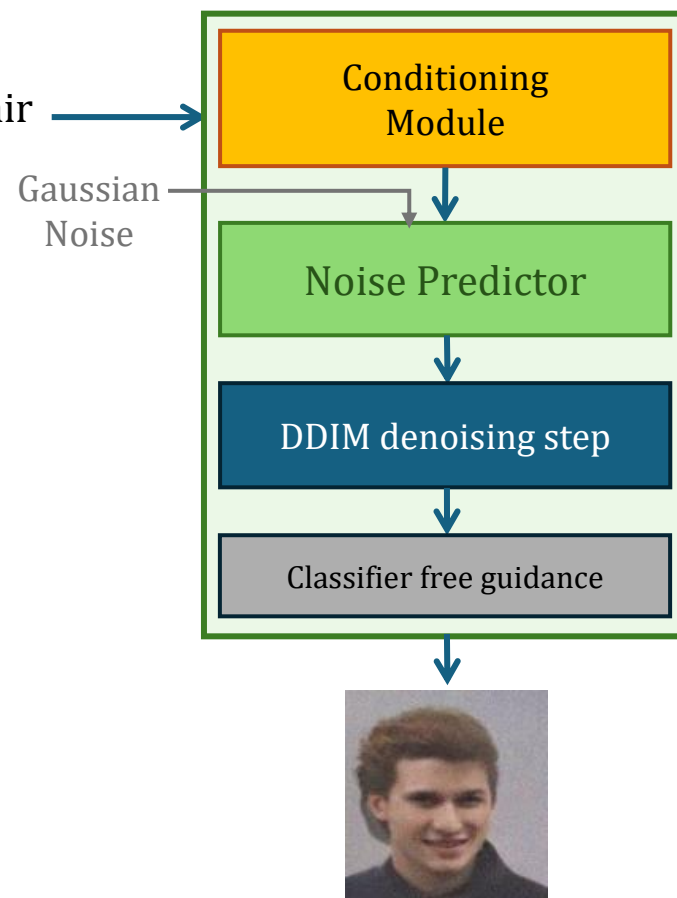
Backgrounds

- **Compositional Conditional Diffusion Models**
 - Capable of generating unseen class images



c_1 : Brown hair
 c_2 : Male

Compositional Conditional Diffusion Model





Motivation

- Can we make it faster (than DDIM) for sampling?

DDIM Reverse Process

20~50 denoising steps!



CFG formulation:

$$(1 + \omega) * x_0^{Cond} - \omega * x_0^{Uncond}$$

Repeat 20~50 times
(inference steps)

40~100 steps for
Conditioned and
Unconditioned outputs

CCDM

Conditioning
Module

Noise Predictor

DDIM denoising step

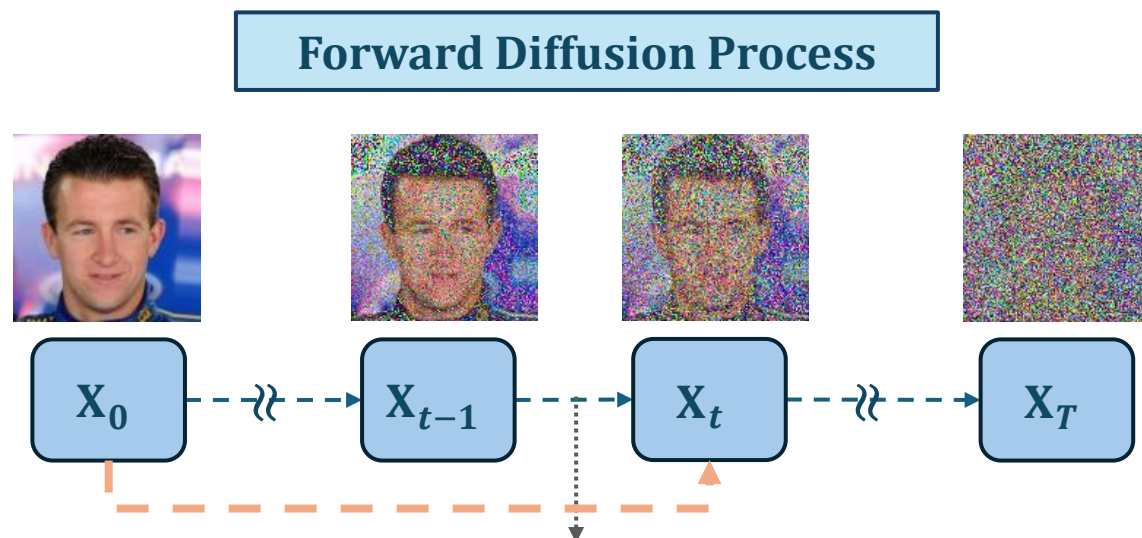
Classifier free guidance

2

Related Works and Preliminaries

Diffusion Denoising Probabilistic Models

- Forward diffusion process and reparameterization



$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} \cdot x_{t-1}, \beta_t \mathbf{I})$$

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s)$$

forward process formula

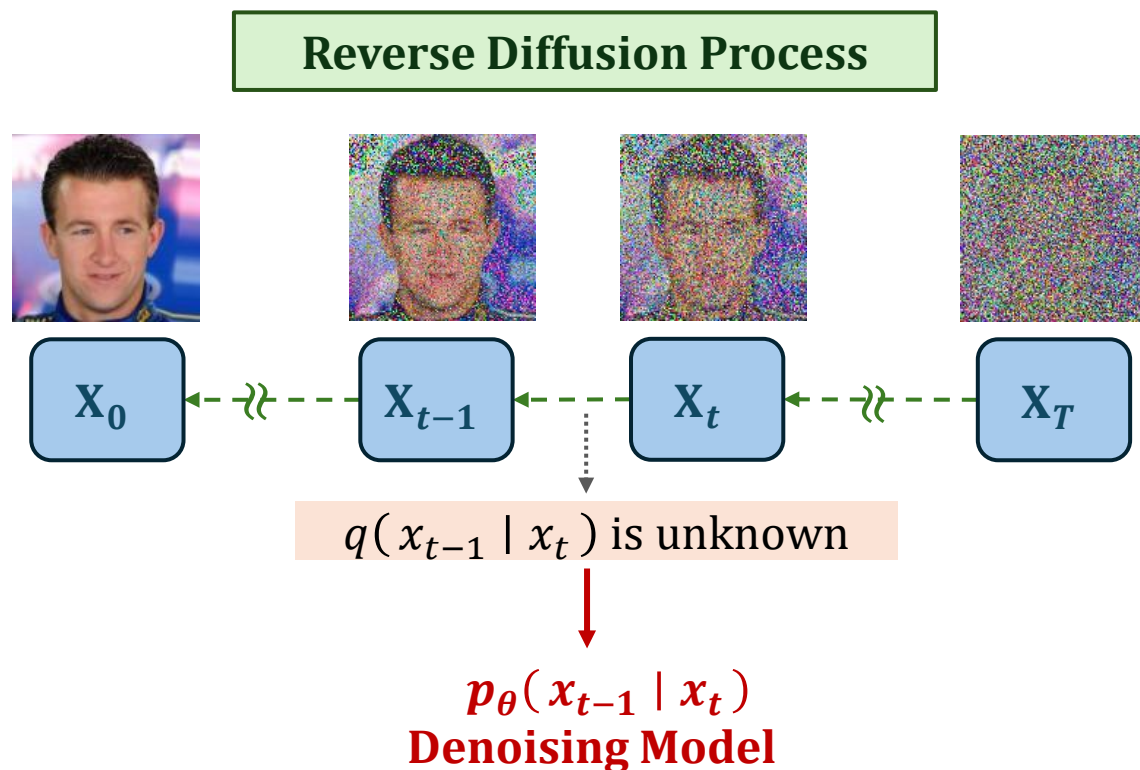
$$x_t = \sqrt{\bar{\alpha}_t} \cdot x_0 + \sqrt{1 - \bar{\alpha}_t} \cdot \epsilon, \quad \epsilon \sim \mathcal{N}(0, \mathbf{I})$$

Reparameterization

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} \cdot x_0, (1 - \bar{\alpha}_t) \cdot \mathbf{I})$$

Diffusion Denoising Probabilistic Models

- Reverse diffusion process, training and sampling



Algorithm 1 Training

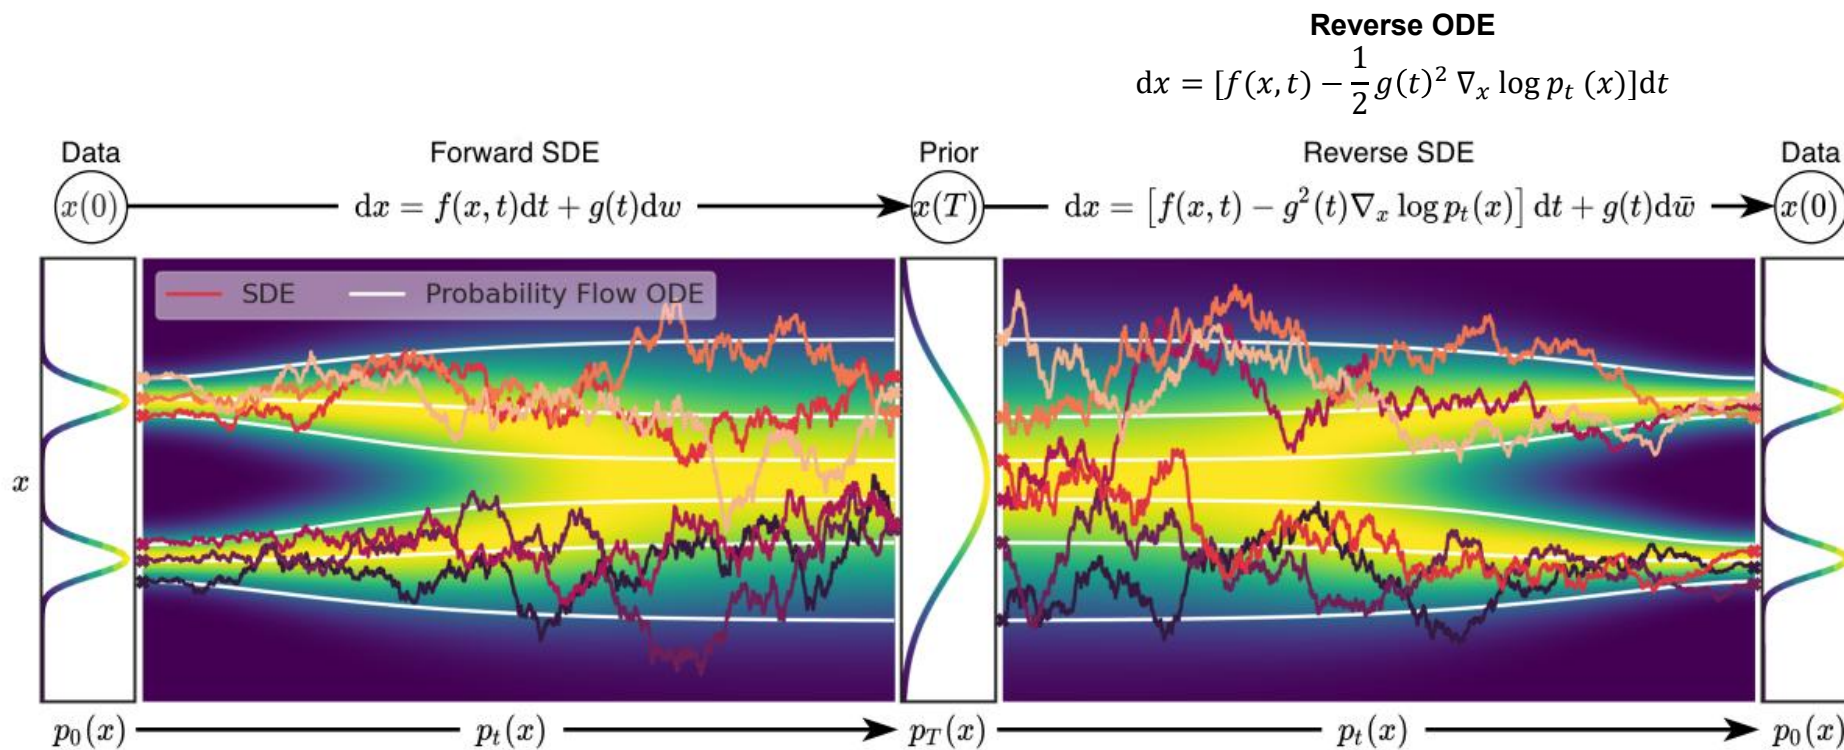
```
1: repeat
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$ 
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take gradient descent step on
MSE loss  $\|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}\mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t)\|^2$ 
6: until converged
```

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_{\theta}(\mathbf{x}_t, t) \right) + \underbrace{\sigma_t \mathbf{z}}_{\text{stochasticity}}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

Score-Based Generative Modeling Through SDE


- Stochastic differential equation & Ordinary differential equation



“For all diffusion processes, there exists a corresponding *deterministic process*, whose trajectories share the same marginal probability densities as the SDE.”

Diffusion Denoising Implicit Models


- Non-Markovian, deterministic denoising
- Speeds up sampling
- No **retraining** required to utilize DDIM sampling



$$\mathbf{x}_{t-1} = \underbrace{\sqrt{\alpha_{t-1}} \left(\frac{\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}{\sqrt{\alpha_t}} \right)}_{\text{"predicted } \mathbf{x}_0"} + \underbrace{\sqrt{1 - \alpha_{t-1} - \sigma_t^2} \cdot \epsilon_{\theta}^{(t)}(\mathbf{x}_t)}_{\text{"direction pointing to } \mathbf{x}_t"} + \underbrace{\sigma_t \epsilon_t}_{\text{random noise}}$$

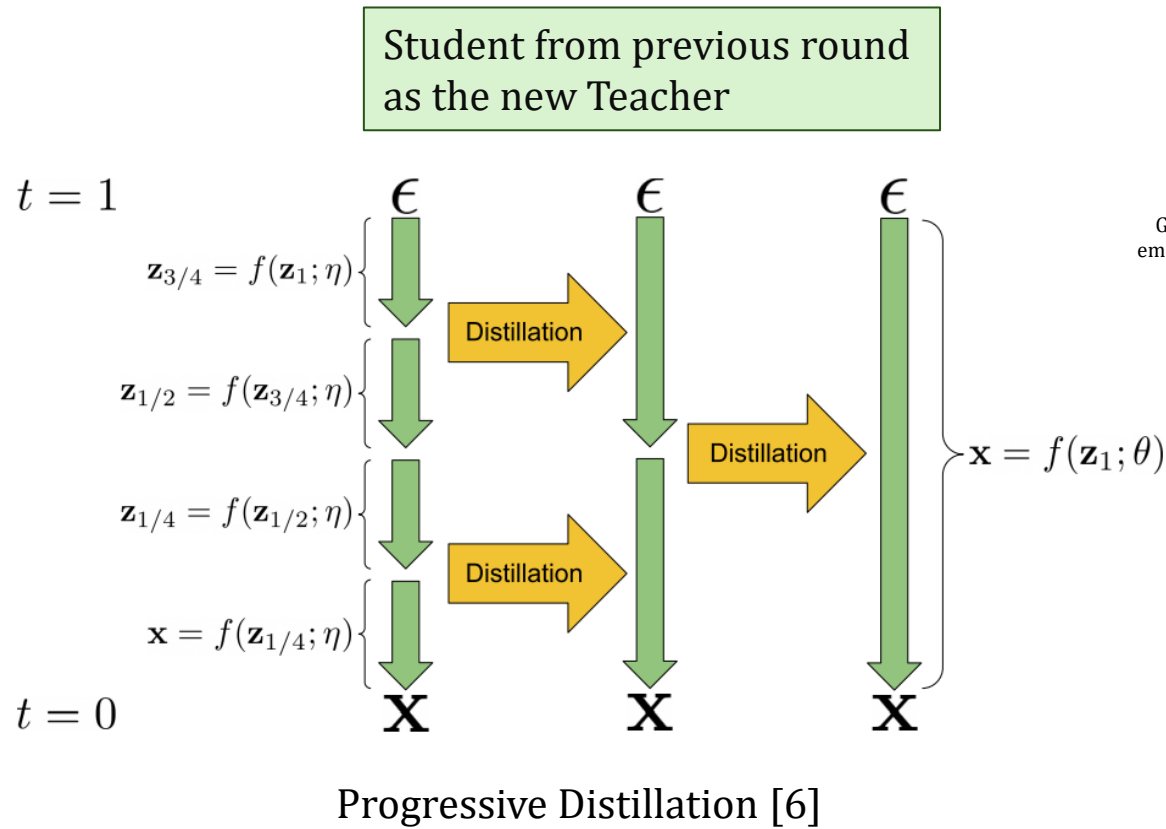
Classifier free guidance

- **Doesn't** need **another classifier** to guide the Diffusion model.
- When training, mostly learns **conditioned** output x^c with input c .
- Occasionally **drops** condition for **unconditioned** output x^\emptyset .
- Weighted combination of conditioned output and unconditioned output.



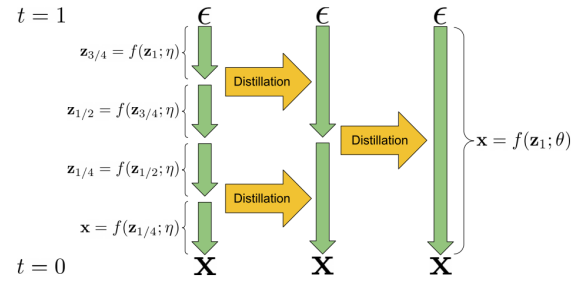
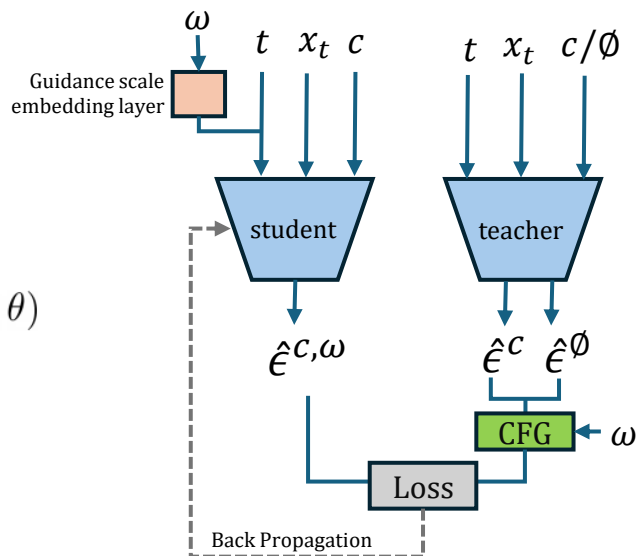
➤ Problem: Needs **twice** (cond+uncond) as **many inference steps**.

Distillation on Diffusion models



Stage 1- eliminate CFG

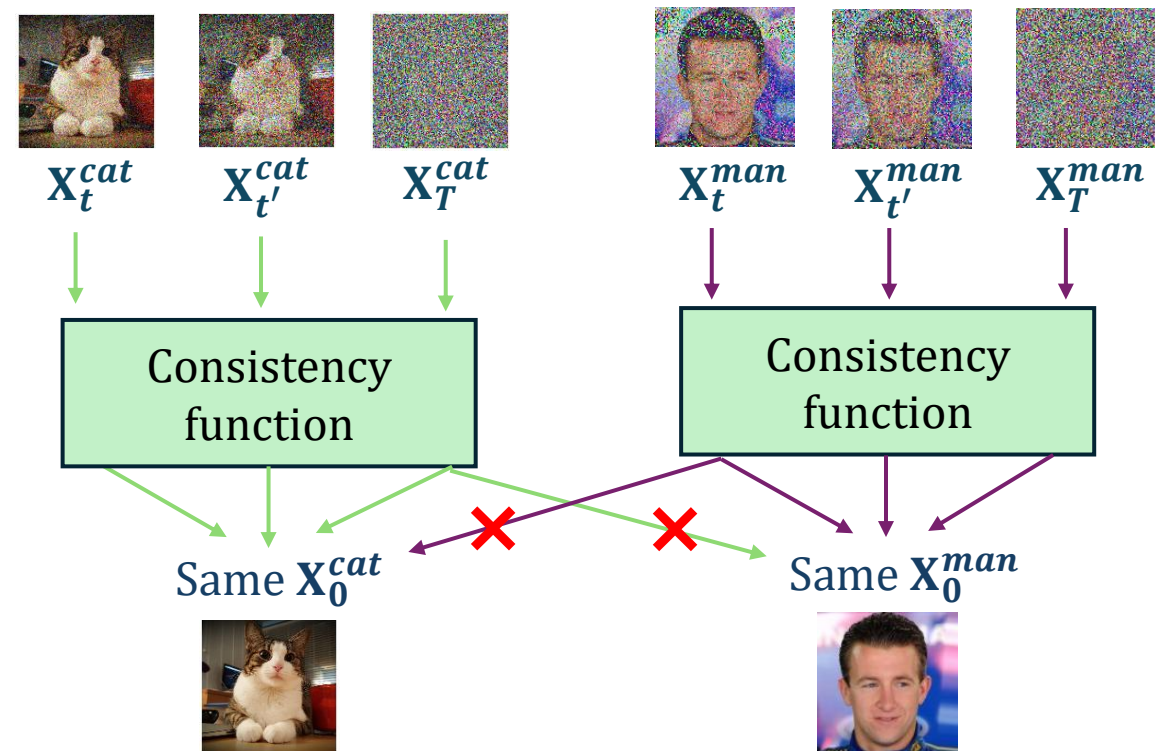
Stage 2- Progressive Distillation



2-stage Distillation on Guided Diffusion models[7]

Consistency Models

- Consistency function
 - Given any $t \in [0, \dots, T]$, exists a function: $f(x_t) \rightarrow x_0$
 - Same x_0 for $x_t, x_{t'}, x_T$ on the same diffusion trajectory



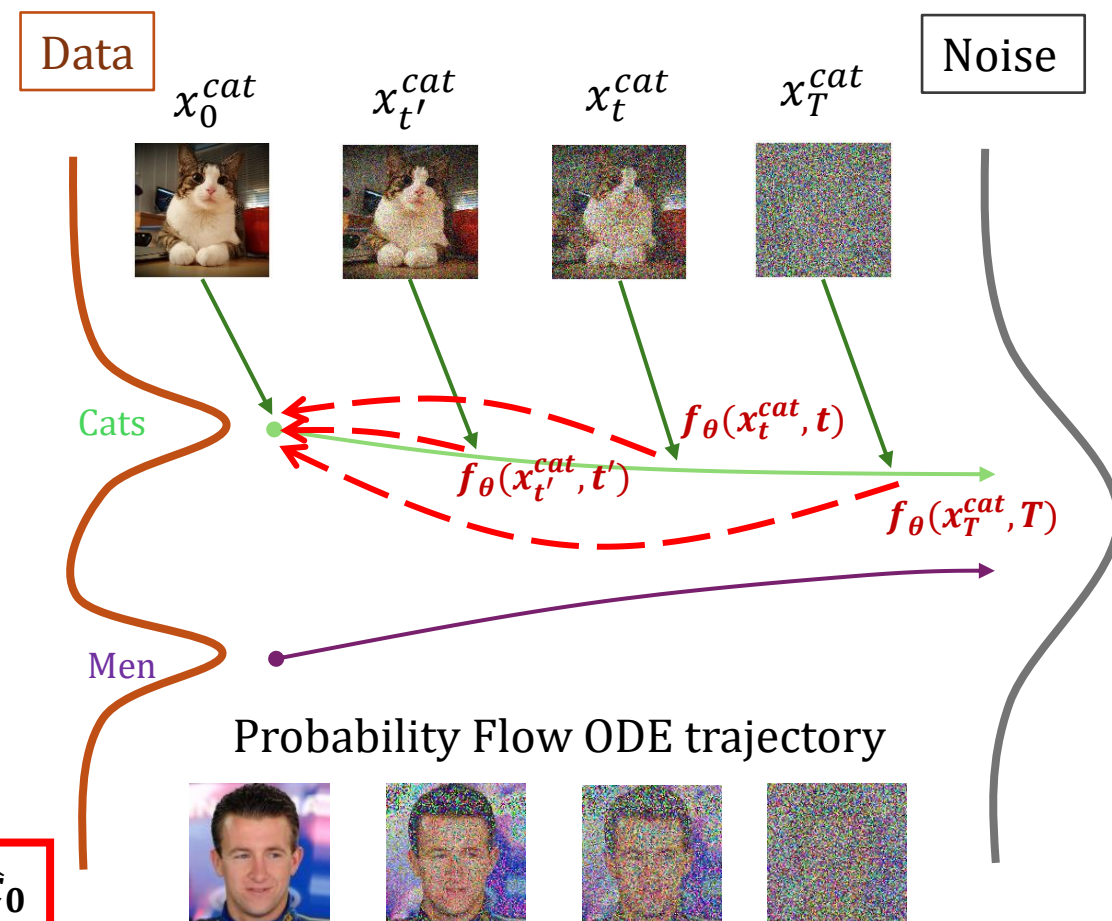
Consistency Models

- Loss function for training
 - Train a **model** f_θ to approximate this function
 - Due to the **deterministic** property of the **PF-ODE**, we may train a consistency model by enforcing the **Consistency objective**
 - Loss function for training:

$$\min_{\theta} [d(f_\theta(x_{t_n}, t_n), f_\theta(x_{t_{n-1}}, t_{n-1}))],$$

So that :

$$f_\theta(x_{t_n}, t_n) = f_\theta(x_{t_{n-1}}, t_{n-1}) = f_\theta(x_{t_1}, t_1) = f_\theta(x_0, t_0) = \hat{x}_0$$



Consistency Models

• Consistency Distillation

- How to obtain adjacent point $x_{t_{n-1}}$?
 - May use forward process formula
 - **Or Better** – Teacher model f_ϕ as the ODE solver Φ

▪ Consistency Distillation

Distill a pretrained Diffusion model into a CM

Algorithm 2 Consistency Distillation (CD)

Input: dataset \mathcal{D} , initial model parameter θ , learning rate η , ODE solver $\Phi(\cdot, \cdot; \phi)$, $d(\cdot, \cdot)$, $\lambda(\cdot)$, and μ

$\theta^- \leftarrow \theta$

repeat

 Sample $\mathbf{x} \sim \mathcal{D}$ and $n \sim \mathcal{U}[1, N - 1]$

 Sample $\mathbf{x}_{t_{n+1}} \sim \mathcal{N}(\mathbf{x}; t_{n+1}^2 \mathbf{I})$

$\hat{\mathbf{x}}_{t_n}^\phi \leftarrow \mathbf{x}_{t_{n+1}} + (t_n - t_{n+1})\Phi(\mathbf{x}_{t_{n+1}}, t_{n+1}; \phi)$

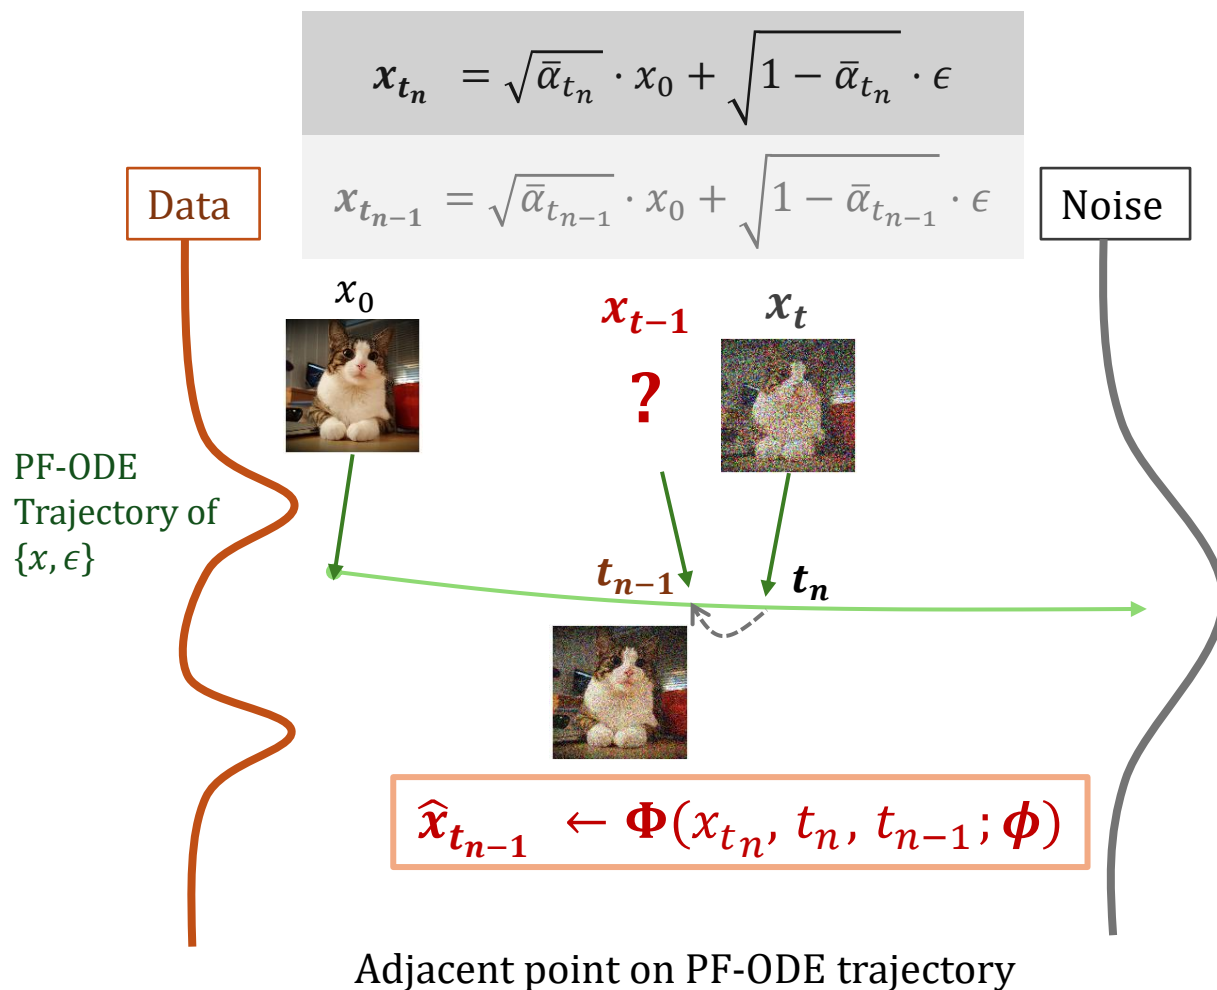
$\mathcal{L}(\theta, \theta^-; \phi) \leftarrow$

$\lambda(t_n)d(f_\theta(\mathbf{x}_{t_{n+1}}, t_{n+1}), f_{\theta^-}(\hat{\mathbf{x}}_{t_n}^\phi, t_n))$

$\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}(\theta, \theta^-; \phi)$

$\theta^- \leftarrow \text{stopgrad}(\mu \theta^- + (1 - \mu)\theta)$

until convergence



Consistency Models

- Sampling

Single step or Multistep Consistency Sampling

Algorithm 1 Multistep Consistency Sampling

Input: Consistency model $f_\theta(\cdot, \cdot)$, sequence of time points $\tau_1 > \tau_2 > \dots > \tau_{N-1}$, initial noise $\hat{\mathbf{x}}_T$

$\mathbf{x} \leftarrow f_\theta(\hat{\mathbf{x}}_T, T)$

for $n = 1$ **to** $N - 1$ **do**

 Sample $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

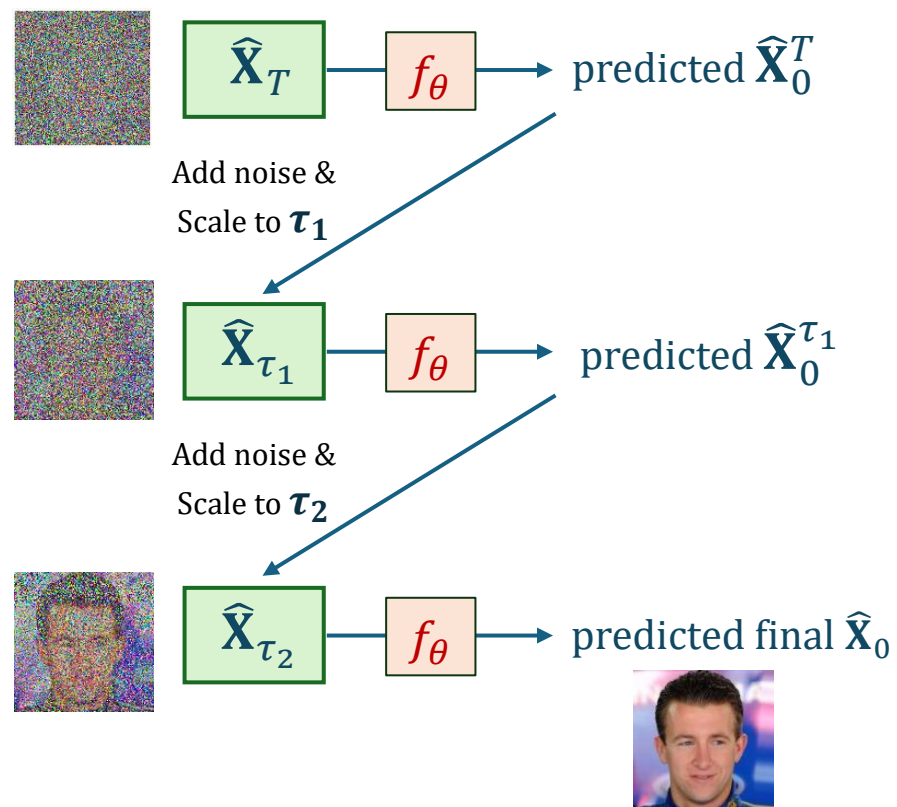
$\hat{\mathbf{x}}_{\tau_n} \leftarrow \mathbf{x} + \sqrt{\tau_n^2 - \epsilon^2} \mathbf{z}$

$\mathbf{x} \leftarrow f_\theta(\hat{\mathbf{x}}_{\tau_n}, \tau_n)$

end for


Output: \mathbf{x}

Example diagram of 3 step sampling

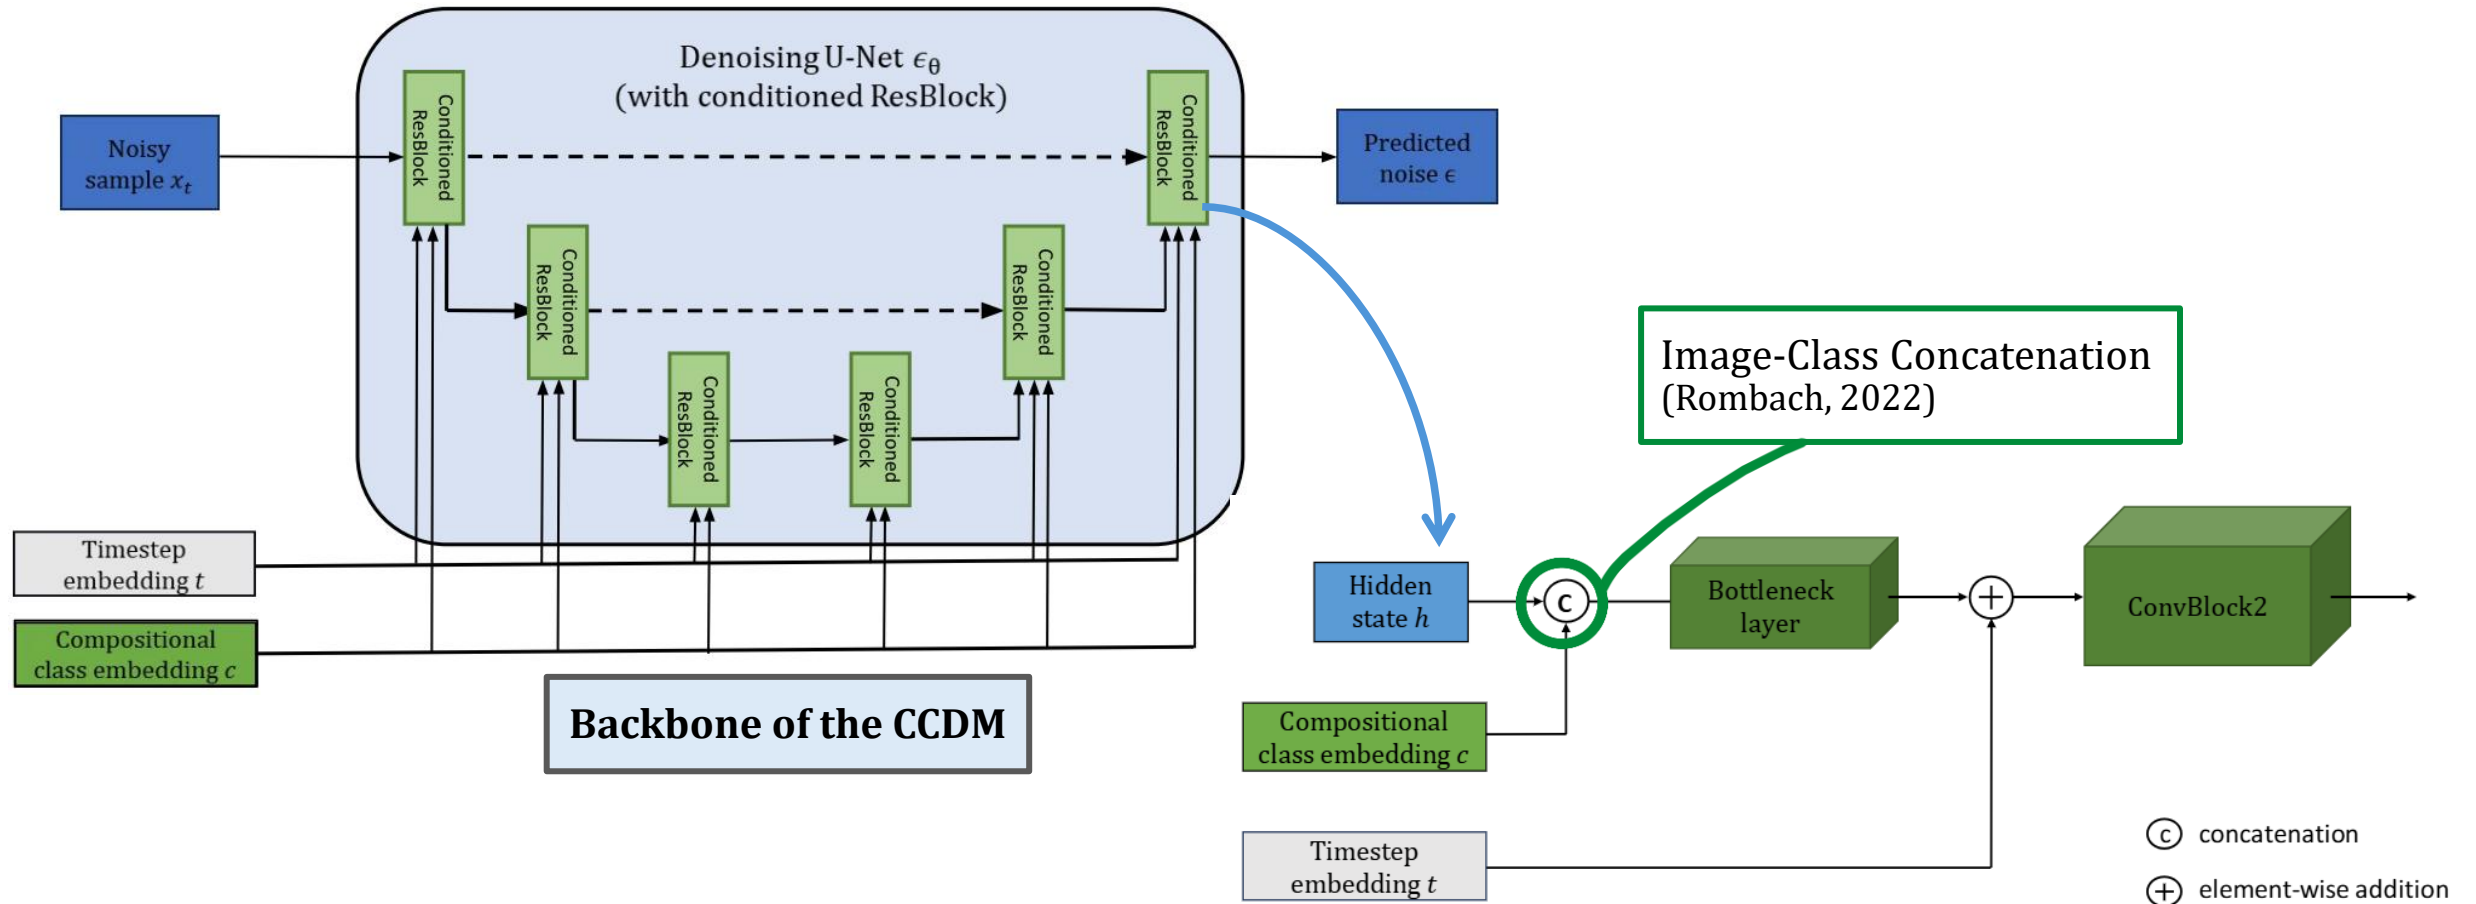


Compositional Zero-Shot Learning

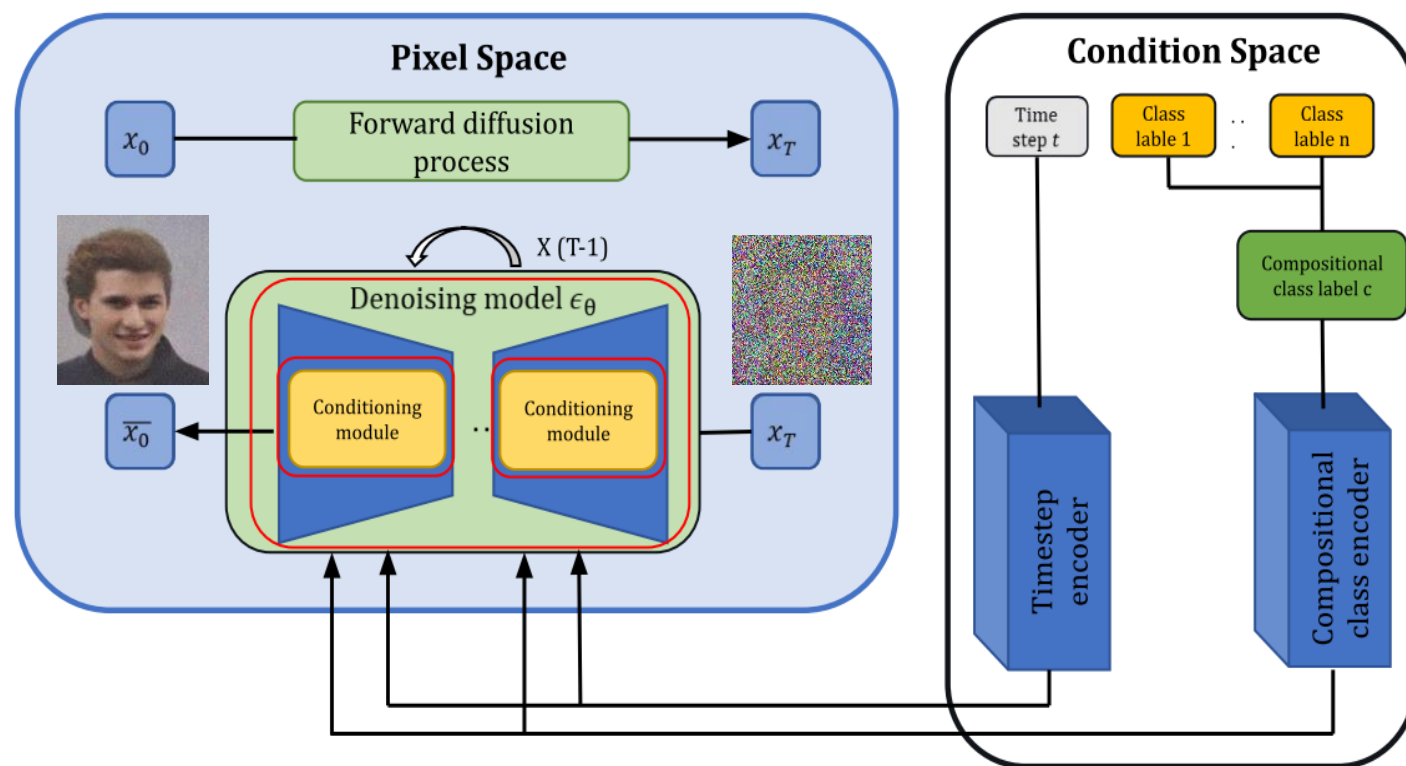
- CZSL allows models to predict unseen classes by leveraging a combination of zero-shot learning and compositional understanding.



Compositional Conditional Diffusion Models



Compositional Conditional Diffusion Models



Algorithm Training CCDM

- 1: **repeat:**
- 2: $(x_0, c) \sim p(x, c)$
- 3: $c \leftarrow \emptyset$ with probability p_{uncond}
- 4: $t \sim Uniform(\{1, \dots, T\})$
- 5: $\epsilon \sim \mathcal{N}(0, \mathbf{I})$
- 6: $x_t = \sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon$
- 7: Take a gradient step on $\nabla_\theta \|\epsilon - \epsilon_\theta(\sqrt{\alpha_t}x_0 + \sqrt{1 - \alpha_t}\epsilon, t, c)\|^2$
- 8: **until** converged

3

Proposed Method

3. Proposed Method

Compositional Conditional Consistency Model

- Using CCDM as the teacher model for consistency distillation

Follows Latent Consistency Model's (Luo, 2023) implementation of:

- skipping step $k=20$
Reduces training time
- guidance scale embedding layer

Notations:

$n \in [1, 2, \dots, T]$: step index

$t_n \in [0, 1]$: time at step n

$\epsilon \sim \mathcal{N}(0, \mathbf{I})$: Gaussian Noise

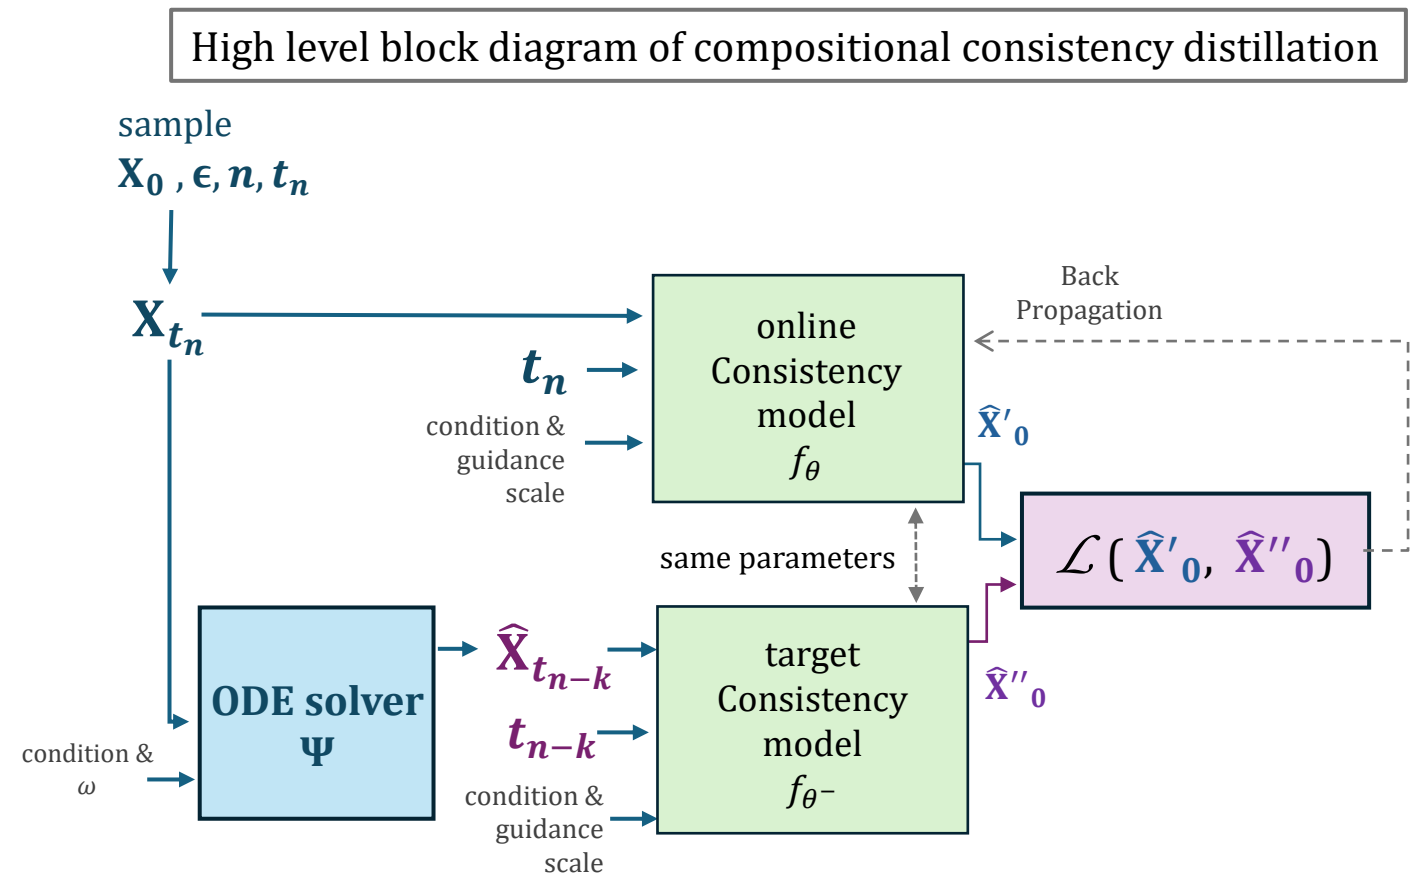
X_{t_n} : noisy image at t_n

\hat{X}_{t_n} : predicted image at t_n

X_0 : clean image from dataset

\hat{X}_0 : predicted clean image

ω : guidance scale



3. Proposed Method

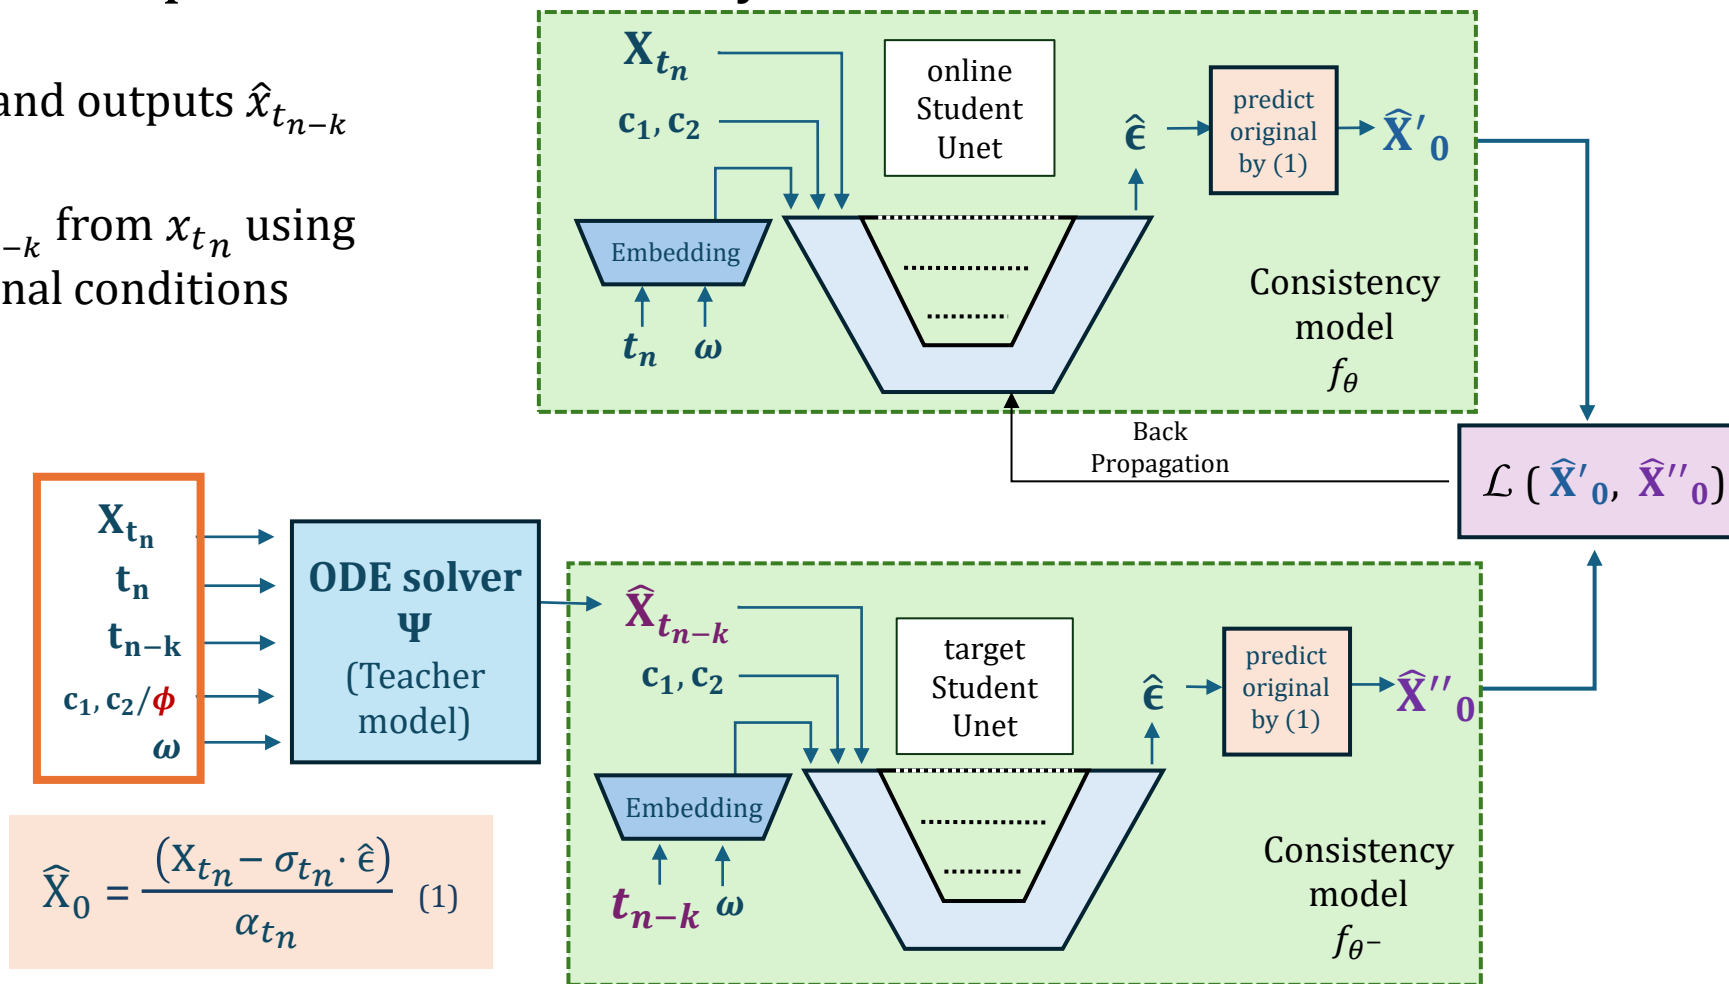
Compositional Conditional Consistency Model

- Detailed block diagram of compositional consistency distillation

- ODE solver takes input x_{t_n} and outputs $\hat{x}_{t_{n-k}}$
- Deterministically estimates $\hat{x}_{t_{n-k}}$ from x_{t_n} using CCDM + DDIM under compositional conditions

Notations:

$n \in [1, 2, \dots, T]$: step index
 $t_n \in [0, 1]$: time at step n
 X_{t_n} : noisy image at t_n
 \hat{X}_{t_n} : predicted image at t_n
 X_0 : clean image from dataset
 \hat{X}_0 : predicted clean image
 α_{t_n} : signal rate at t_n
 σ_{t_n} : noise rate at t_n
 $\hat{\epsilon}$: predicted noise
 ω : guidance scale



3. Proposed Method

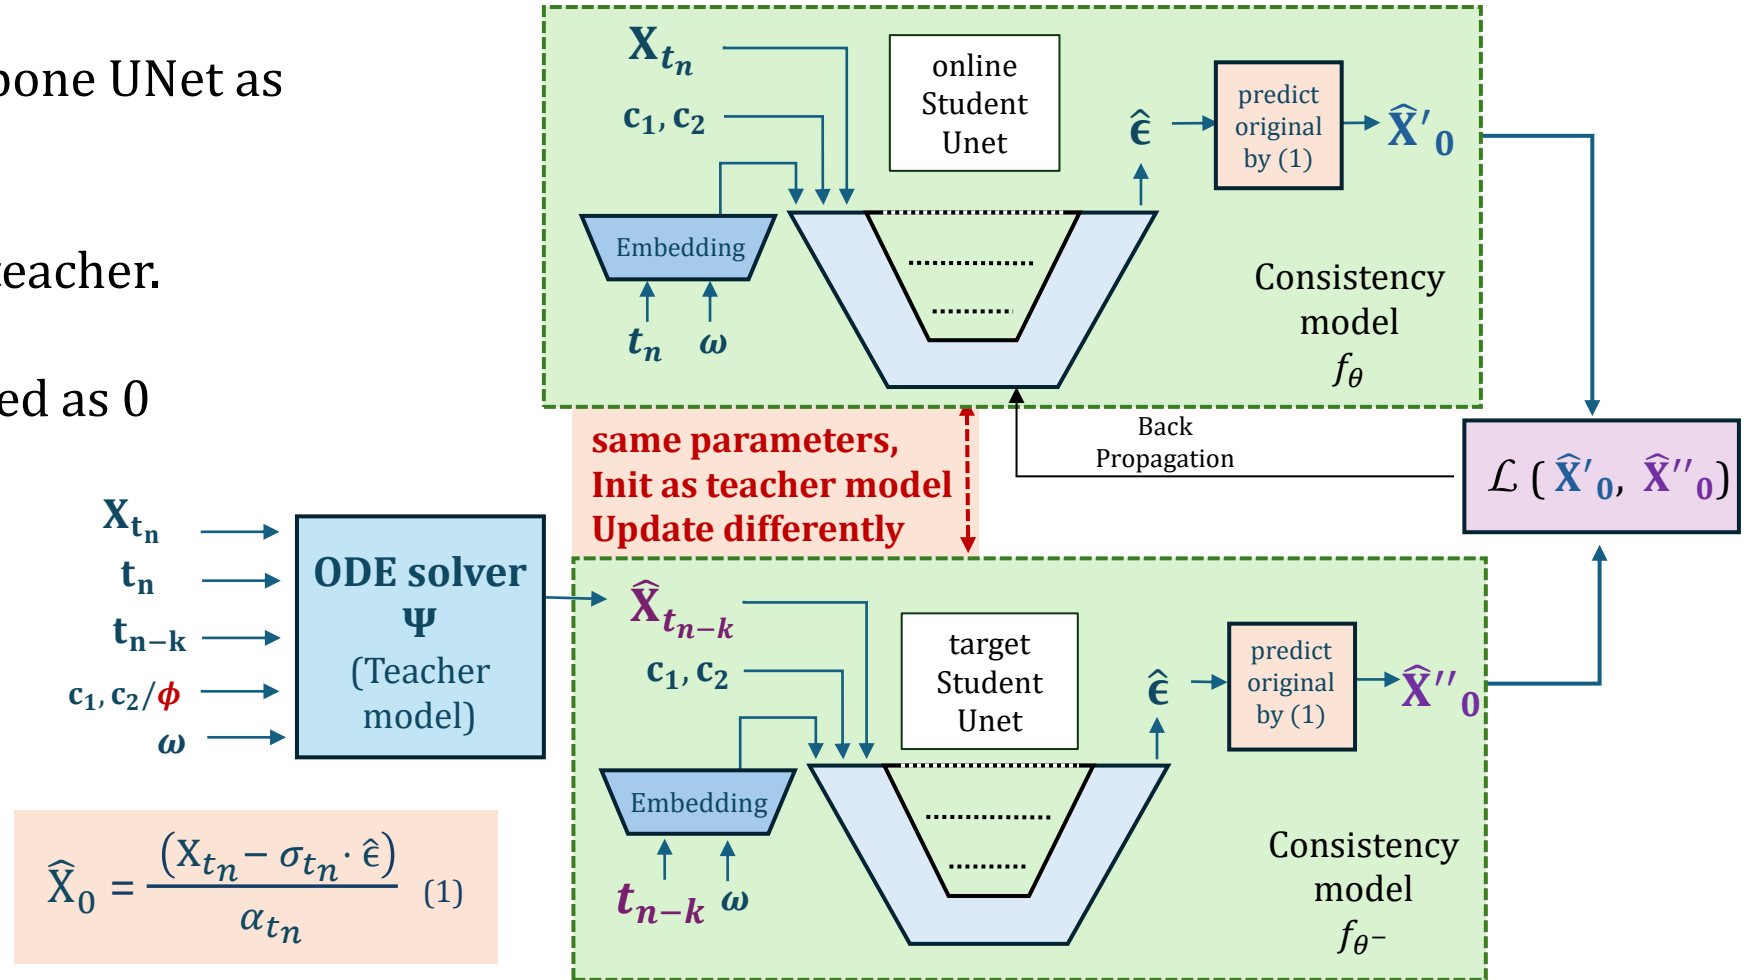
Compositional Conditional Consistency Model

- Detailed block diagram of compositional consistency distillation

- f_θ and f_{θ^-} using same backbone UNet as teacher (CCDM)
- both weights initialized as teacher.
- ω embedding layer initialized as 0

Notations:

$n \in [1, 2, \dots, T]$: step index
 $t_n \in [0, 1]$: time at step n
 X_{t_n} : noisy image at t_n
 \hat{X}_{t_n} : predicted image at t_n
 X_0 : clean image from dataset
 \hat{X}_0 : predicted clean image
 α_{t_n} : signal rate at t_n
 σ_{t_n} : noise rate at t_n
 $\hat{\epsilon}$: predicted noise
 ω : guidance scale



3. Proposed Method


Compositional Conditional Consistency Model

- Detailed block diagram of compositional consistency distillation

- Target UNet updated from online UNet using Exponential Moving Average (EMA)

Notations:

$n \in [1, 2, \dots, T]$: step index
 $t_n \in [0, 1]$: time at step n
 X_{t_n} : noisy image at t_n
 \hat{X}_{t_n} : predicted image at t_n
 X_0 : clean image from dataset
 \hat{X}_0 : predicted clean image
 α_{t_n} : signal rate at t_n
 σ_{t_n} : noise rate at t_n
 $\hat{\epsilon}$: predicted noise
 ω : guidance scale



3. Proposed Method

Ways of Computing $X_{t_{n-k}}$

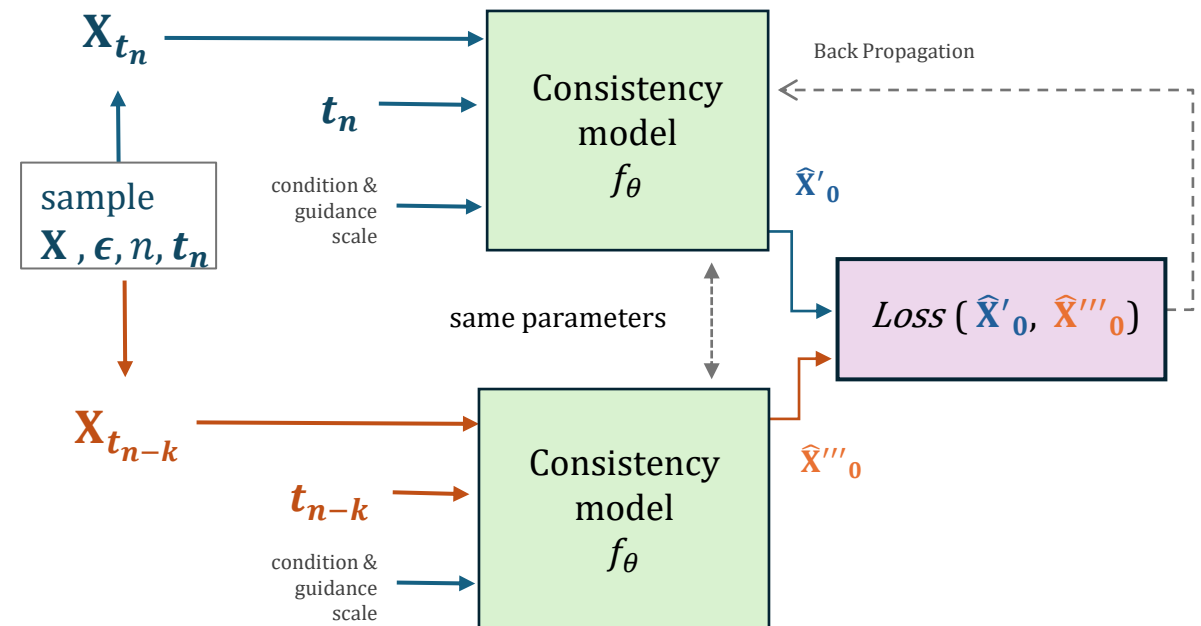
- Teacher predicted or Forward-process formulation?

- CCCM with teacher's supervision offers high-quality samples.

- Could formulated $x_{t_{n-k}}$ do better?
 - Not quite if solely rely on formulated.

$$x_{t_{n-k}} = \sqrt{\bar{\alpha}_{t_{n-k}}} \cdot x_0 + \sqrt{1 - \bar{\alpha}_{t_{n-k}}} \cdot \epsilon, \epsilon \text{ same as in } x_{t_n}$$

- Gradually shifting from teacher to formulated $x_{t_{n-k}}$ might help?



3. Proposed Method

Epoch function $\lambda(e)$ - Fusion implementation

- To control the strength of 2 branches: $\hat{x}_{t_{n-k}}$ and $x_{t_{n-k}}$
 - Epoch function $\lambda(e): \{1, 2, \dots, e_{max}\} \rightarrow [0, 1]$.
1 = fully teacher signal, 0 = formulated signal
 - Switch, Step Fuse, and Loss Fuse implemented via $\lambda(e)$
 - Controlled by Fuse Scheduler

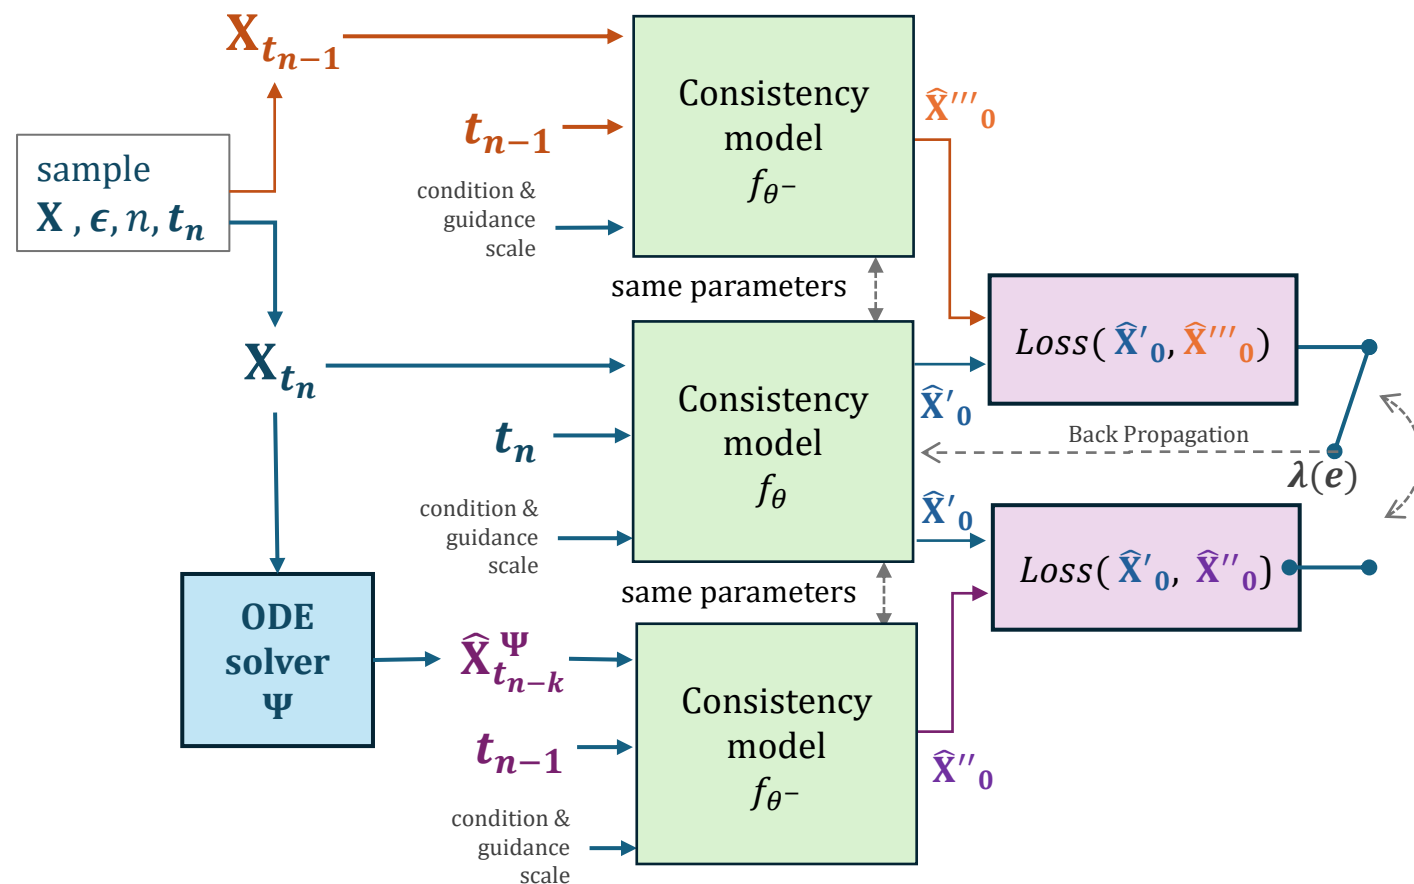
3. Proposed Method

Modified Consistency Distillation

- Switch Strategy

- Switches source of $x_{t_{n-k}}$ based on an epoch threshold :

$$\lambda(e) = \begin{cases} 1, & e < \text{threshold} \\ 0, & e \geq \text{threshold} \end{cases}$$



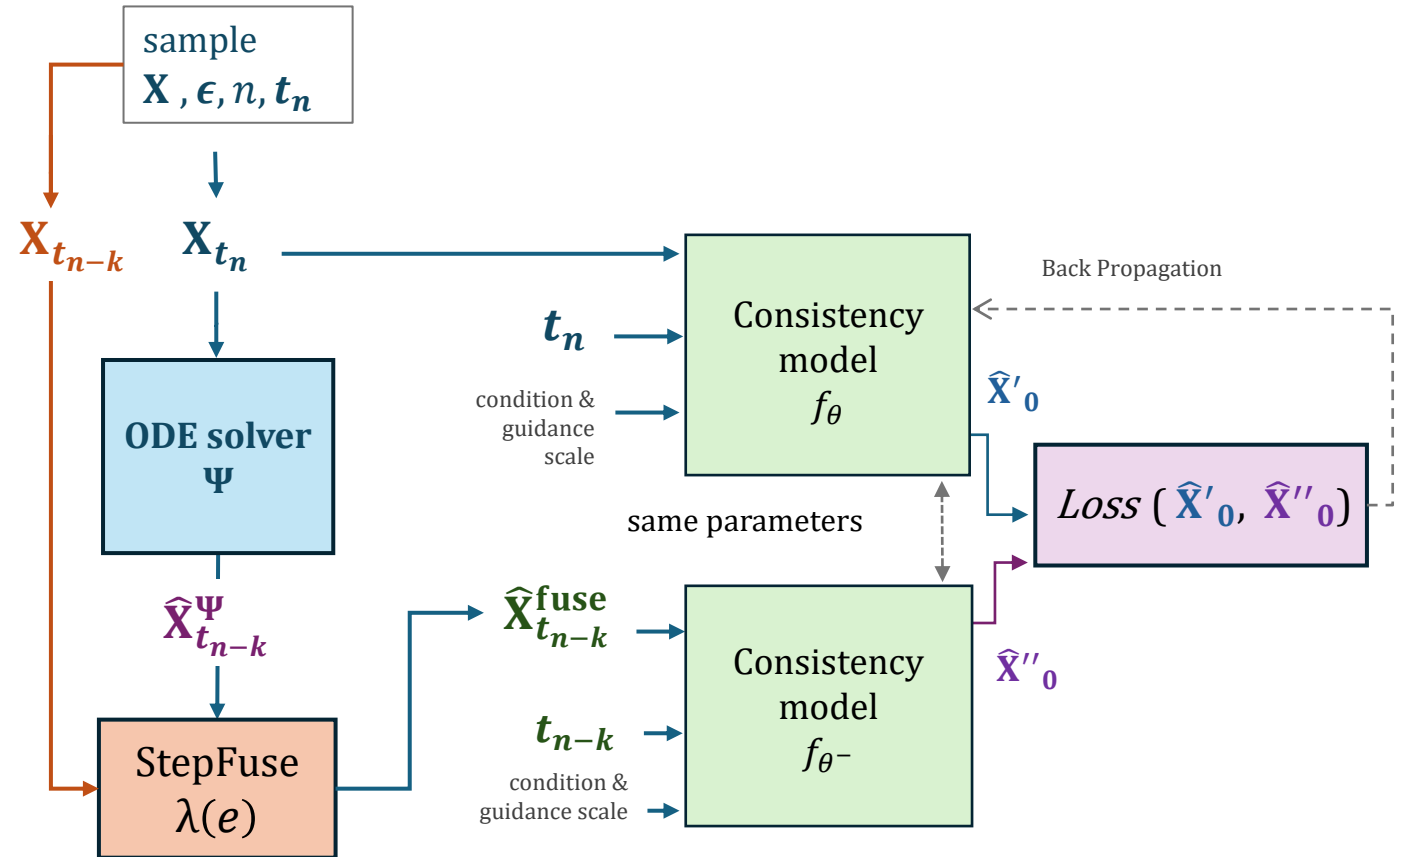
3. Proposed Method

Modified Consistency Distillation

- Step-Fuse Strategy

- Pixel-wise weighted sum of two x_{t-k} sources:

$$\hat{x}_{t_{n-k}}^{\text{fuse}} = \lambda(e) \cdot \hat{x}_{t_{n-k}}^{\Psi} + (1 - \lambda(e)) \cdot x_{t-k}$$



3. Proposed Method

Modified Consistency Distillation

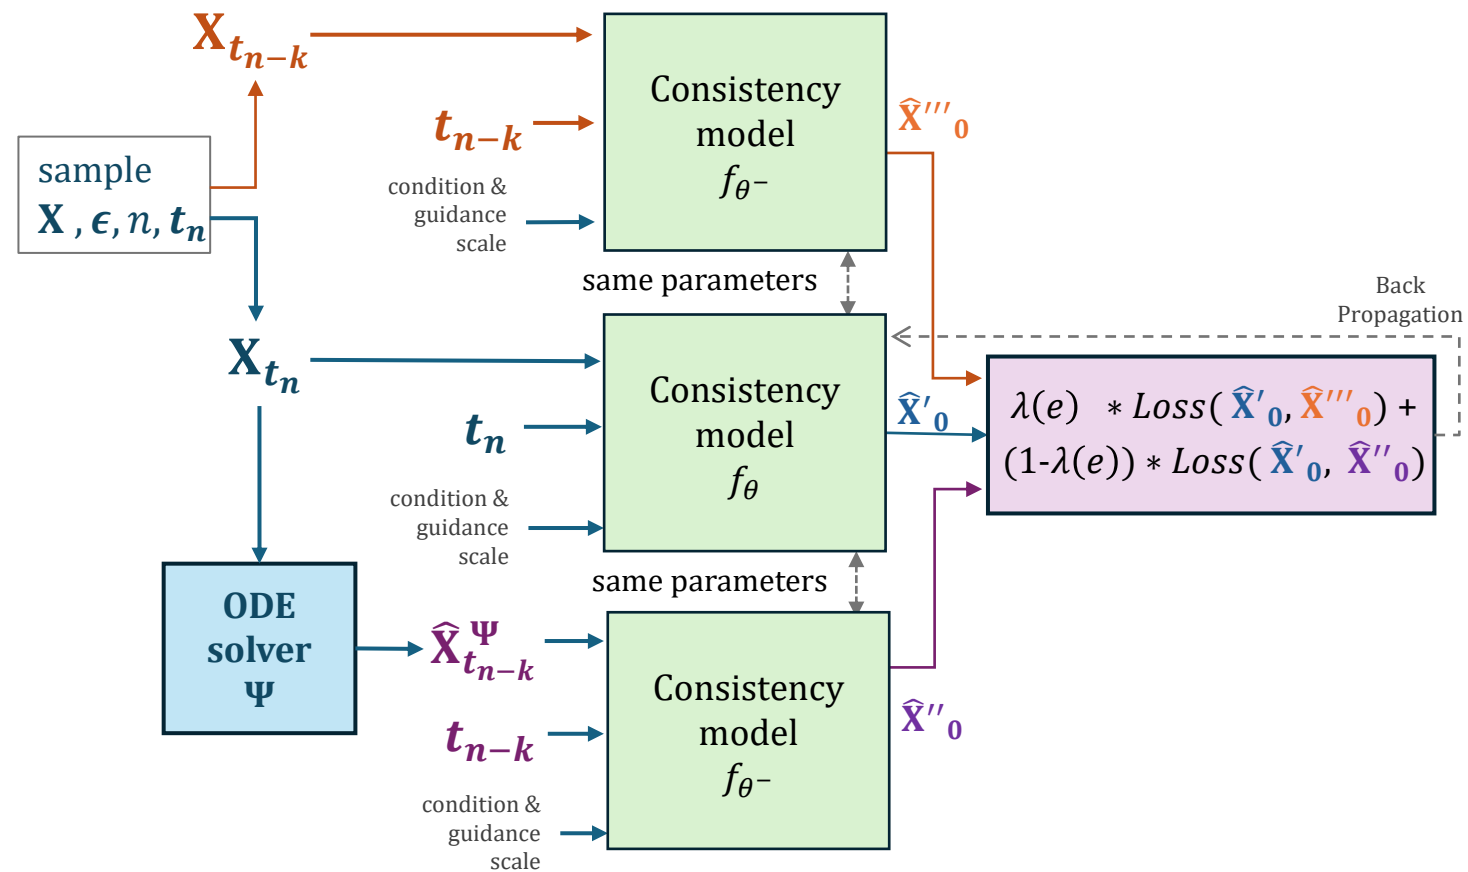
- Loss-Fuse Strategy

- Weighted sum of two loss terms:

$$\mathcal{L}_{teacher} = \mathbb{E} \left[d \left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta}(\hat{x}_{t_{n-k}}^{\Psi}, t_{n-k}, c_1, c_2, \omega) \right) \right]$$

$$\mathcal{L}_{formulated} = \mathbb{E} \left[d \left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta}(x_{t_{n-k}}, t_{n-k}, c_1, c_2, \omega) \right) \right]$$

$$loss = \lambda(e) \cdot \mathcal{L}_{teacher} + (1 - \lambda(e)) \cdot \mathcal{L}_{formulated}$$



3. Proposed Method

Fuse Scheduler design

- Constant:

$$\lambda(epoch) = \lambda_0 ,$$

fixed blending weight throughout training.

- $epoch \in [1, epoch_{max}]$

- $prog = \frac{epoch}{epoch_{Max}} \in [0,1]$

- Exponential Decay:

$$\lambda(epoch) = e^{-\gamma \cdot prog} \cdot (1 - prog),$$

γ = decay rate

- Piecewise Linear:

$$\lambda(epoch) = \lambda_i + \frac{\lambda_{i+1} - \lambda_i}{p_{i+1} - p_i} (prog - p_i)$$

Defined by control points $\{(p_i, \lambda_i)\}$, $p_i, \lambda_i \in [0,1]$.

Linearly interpolated over epochs

3. Proposed Method

Summary of proposed Strategies

Method	Source of x_{t-k}	Description
Fully Teacher	Teacher Model Prediction (via ODE Solver)	Uses noise predicted by teacher model to estimate x_{t-k} through an ODE solver.
Fully diffusion formula	Reused Forward Process Noise	x_{t-k} formulated using the same noise as x_t .
Step Fusion	Mixed (Teacher & Diffusion Formula)	Pixel-wise weighted combination of both teacher and formulated x_{t-k} .
Loss Fusion	Both (Teacher & Diffusion Formula)	Computes separate losses, combines with weighting.
Switch	Alternating	Uses teacher before an epoch threshold, switches to formulated x_{t-k} after.

4

Experiments

4.Experiments

Experiment setup

- Dataset preparation
 - c_1 : Hair color, 4 classes
 - c_2 : Gender, 2 classes
 - Total of $4*2=8$ composition classes,
 - Unseen: (Brown hair, Male)
 - Dataset image counts: 1k~20k
 - Image size: 128*128

		c_1			
		Gray	Black	Blonde	Brown
c_2	Male				
	Female				

Training set with compositional class labels c_1, c_2

4.Experiments

Experiment setup




- CCCM and Baseline Training Configuration

Baseline CDDM

Training Epochs	120
Learning rate, scheduler	5e-5 linear with warmup
Loss type	L2-norm

Our CCCM

Training Epochs	80
Learning rate, scheduler	5e-6 constant with warmup
Loss type	Huber loss
Guidance scale interval	2.6, 3.0
Fuse Method	Fully teacher



Experiment setup

- Modified Consistency Distillation Configuration

Fully Teacher

Fuse Method	Fully Teacher
Fuse Scheduler	Constant = 1

Fully Formulated x_{t-k}

Fuse Method	Fully Formulated
Fuse Scheduler	Constant = 0

Switch 32

Fuse Method	Switch
Fuse Scheduler	Threshold = 32

Switch 48

Fuse Method	Switch
Fuse Scheduler	Threshold = 48

Unchanged Hyperparameters

Training Epochs	80
Learning rate, scheduler	5e-6 constant with warmup
Loss type	Huber loss
Guidance scale interval	2.6, 3.0

Experiment setup

- Modified Consistency Distillation Configuration

Loss Fuse Constant

Fuse Method	Loss Fuse
Fuse Scheduler	Constant = 0.8

Loss Fuse piecewise

Fuse Method	Loss Fuse
Fuse Scheduler	Piecewise Linear, (40, 0.5)

$\lambda(epoch)$ drops to 0.5 at epoch 40 linearly, holds at 0.5 till finished.

Loss Fuse exponential

Fuse Method	Loss Fuse
Fuse Scheduler	Exponential Decay, $\gamma = 2.0$

Step Fuse exponential


Fuse Method	Step Fuse
Fuse Scheduler	Exponential Decay, $\gamma = 2.0$

Unchanged Hyperparameters

Training Epochs	80
Learning rate, scheduler	5e-6 constant with warmup
Loss type	Huber loss
Guidance scale interval	2.6, 3.0

Experiment setup

- Evaluation metrics
 - **FID score:**
measures the similarity between generated and real image distributions.
- Each model generates ~7k images (500 or 1000 per class) for evaluation
- **Unseen Class Accuracy** by human evaluation:
Compositional class (*Brown_hair, Male*) is generated, to test the compositional zero-shot image generation ability.
- Each model generates 300 (*Brown_hair, Male*) images.

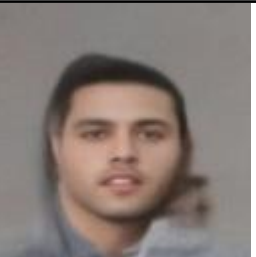
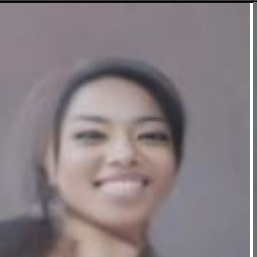


Standards:

- **Unseen class judged by both attributes.**
Masculine features required (e.g., short hair, no makeup).
Feminine traits → failure, even if hair color is correct.

















Qualitative result

- Baseline CCDM vs Fully teacher CCCM

4-step sampling	Blonde Hair		Black Hair		Brown Hair		Gray Hair	
	Male	Female	Male	Female	Male	Female	Male	Female
Baseline CCDM (with 8 inferences)								
								

Qualitative result

- Baseline CCDM vs Fully teacher CCCM

2-step sampling	Blonde Hair		Black Hair		Brown Hair		Gray Hair	
	Male	Female	Male	Female	Male	Female	Male	Female
Baseline CCDM (with 4 inferences)								
CCCM (with 2 inferences)								

4.Experiments

Quantitative result

- FID scores under 2,3 and 4 steps sampling

2 steps

Method	FID score ↓
Baseline DDIM	207.99
Forward-process x_{t-k}	115.66
Step Fuse (exponential)	97.13
Switch (threshold = 32)	94.99
Switch (threshold = 48)	93.97
Loss Fuse (exponential)	91.27
Loss Fuse (piecewise = 40:0.5)	88.82
Loss Fuse (constant = 0.8)	88.47
Fully Teacher x_{t-k}	81.30

3 steps

Method	FID score ↓
Baseline DDIM	136.68
Forward-process x_{t-1}	92.15
Step Fuse (exponential)	83.73
Switch (threshold = 32)	85.66
Switch (threshold = 48)	84.23
Loss Fuse (exponential)	82.90
Loss Fuse (piecewise = 40:0.5)	80.16
Loss Fuse (constant = 0.8)	77.85
Fully Teacher x_{t-k}	73.27

4 steps

Method	FID score ↓
Baseline DDIM	94.67
Diffusion-formulated x_{t-1}	80.94
Step Fuse (exponential)	77.27
Switch (threshold = 32)	76.26
Switch (threshold = 48)	75.86
Loss Fuse (exponential)	75.34
Loss Fuse (piecewise = 40:0.5)	73.45
Loss Fuse (constant = 0.8)	70.15
Fully Teacher x_{t-1}	68.11

4.Experiments

Unseen Accuracy evaluation

- 300 images of Brown hair Male, compositional zero shot generation
4 steps sampling

Method	Acc% ↑
Baseline DDIM	40.6%
Diffusion-formulated $x_{t_{n-k}}$	43.0%
Step Fuse (exponential)	51.6%
Switch (threshold = 32)	49.3%
Switch (threshold = 48)	49.6%
Loss Fuse (exponential)	51.6%
Loss Fuse (piecewise = (40, 0.5))	52.0%
Loss Fuse (constant = 0.8)	50.6%
Fully Teacher $x_{t_{n-k}}$	47.0%



Loss Fuse
piecewise (40,05)



Fully teacher $x_{t_{n-k}}$

5.

Conclusions

Conclusion

- We propose the **Compositional Conditional Consistency Model**.
- Achieving **faster sampling** speed than CCDM.
- **Preserves unseen** class generation.
- Observed that modified consistency distillation strategies yield **better unseen accuracy**.

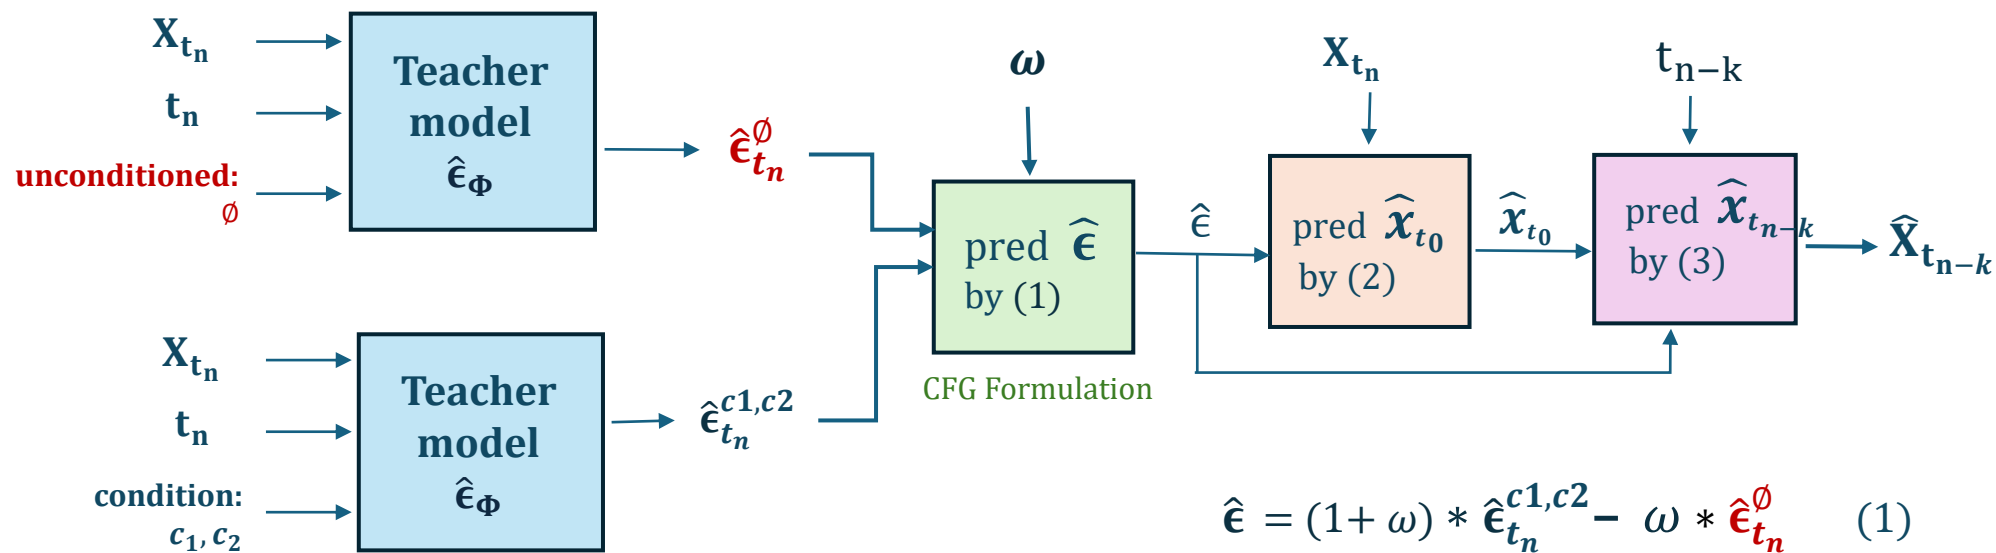
Our guess is: the teacher model may introduce **bias**, encouraging the student to generate seen or **high-confidence images**, which **limits generalization**.

References

- **"Compositional Conditional Diffusion Model"**, S. -L. Lai, P. -C. Chen & C. -W. Ma, 2024.
- **"Consistency Models"**, Y. Song, P. Dhariwal, M. Chen & I. Sutskever, 2023.
- **"Latent consistency models: Synthesizing high-resolution images with few-step inference"**, Luo, S., Tan, Y., Huang, L., Li, J., & Zhao, H. 2023.
- **"Improved techniques for training consistency models"**, Y. Song and P. Dhariwal, 2023.
- **"Score-based generative modeling through stochastic differential equations"**, Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon & B. Poole, 2020.
- **"Denoising diffusion implicit models"**, J. Song, C. Meng & S. Ermon, 2020.
- **"Progressive distillation for fast sampling of diffusion models"**, T. Salimans and J. Ho, 2022.
- **"On distillation of guided diffusion models"**, C. Meng, R. Rombach, R. Gao, D. Kingma, S. Ermon, J. Ho, and T. Salimans, 2023.
- **"Classifier-free diffusion guidance"**, J. Ho and T. Salimans, 2022.
- **"High-resolution image synthesis with latent diffusion models"**, R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer, 2022.

ODE solver

- CCCM uses DDIM as ODE solver



Notations:

t_n : time at step n , $t_n \in [0, 1], n \in [1..T]$

α_{t_n} : signal rate ; σ_{t_n} : noise rate

X_{t_n} : image at t_n ; \hat{X}_{t_n} : predicted image at t_n


$\hat{\epsilon}_{t_n}$: predicted noise at t_n

ω : guidance scale

$$\hat{X}_{t_0} = \frac{X_{t_n} - \sigma_{t_n} * \hat{\epsilon}}{\alpha_{t_n}} \quad (2)$$

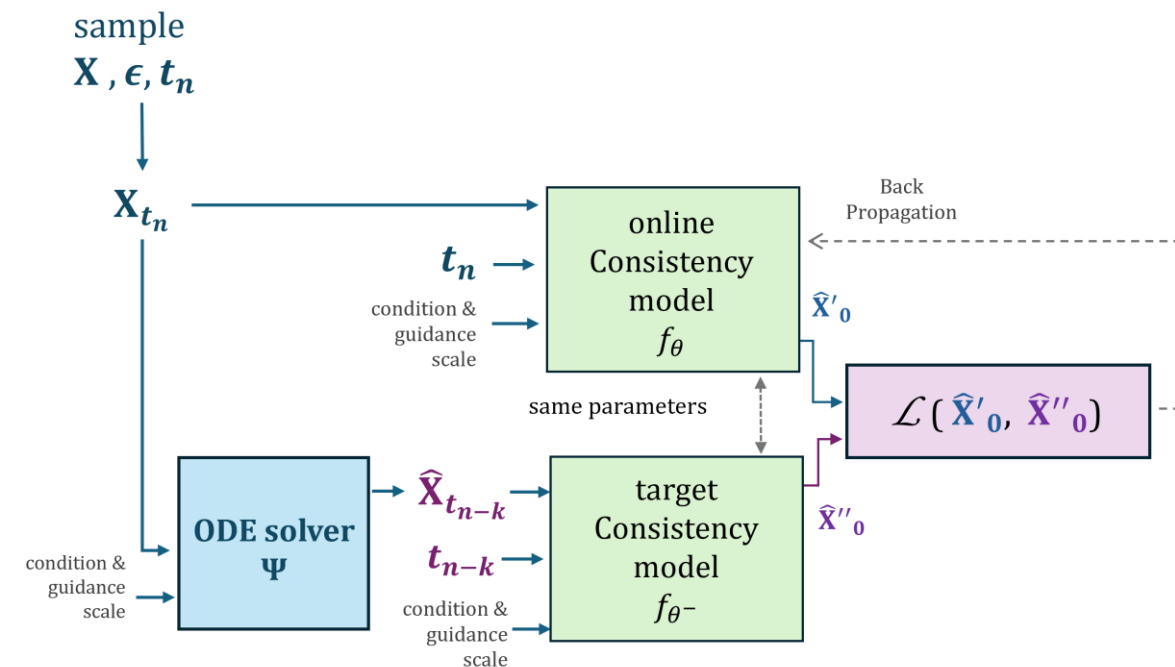
$$\hat{X}_{t_{n-k}} = \alpha_{t_{n-k}} * \hat{X}_0 + \sigma_{t_{n-k}} * \hat{\epsilon} \quad (3)$$

Guidance scale embedding



Compositional Consistency Distillation Algo

- Fully teacher



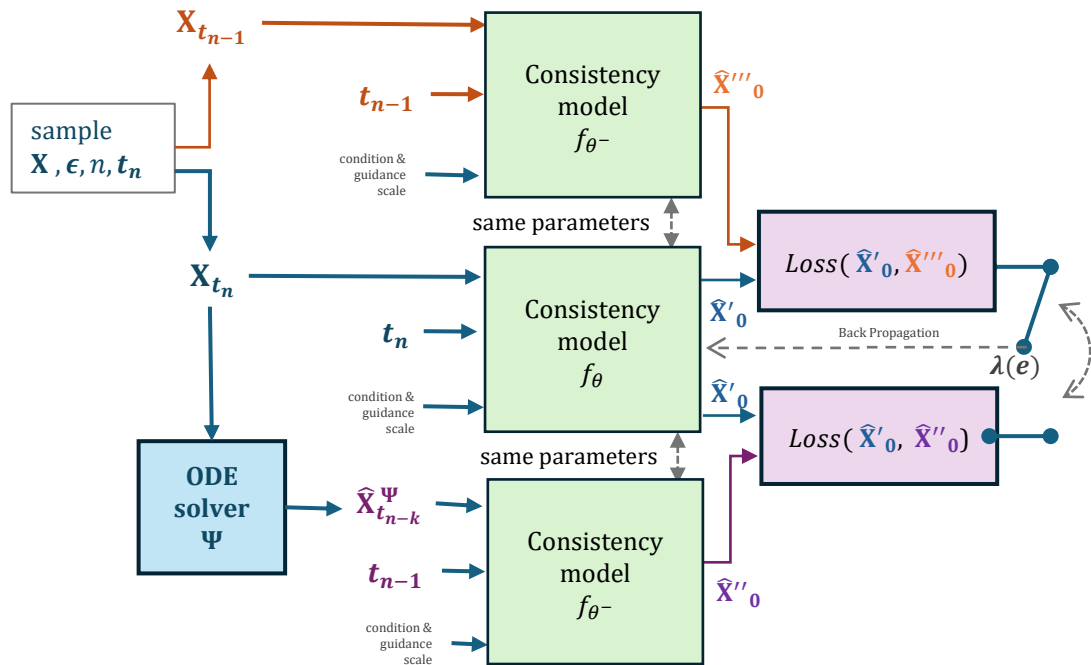
Algorithm 4.1 Compositional label Consistency Distillation

Input: dataset \mathcal{D} , initial model parameter θ , learning rate η , ODE solver $\Psi(\cdot, \cdot, \cdot, \cdot)$, distance metric $d(\cdot, \cdot)$, guidance scale $[\omega_{min}, \omega_{max}]$, skipping steps k

- 1: $\theta^- \leftarrow \theta$
- 2: **repeat**
- 3: Sample $(x, c_1, c_2) \sim \mathcal{D}$, and $\omega \sim [\omega_{min}, \omega_{max}]$
- 4: Sample $\epsilon \sim \mathcal{N}(0, I)$, $n \sim \mathcal{U}[1 + k, N]$
- 5: $x_{t_n} \leftarrow \sqrt{\bar{\alpha}_{t_n}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_n}} \cdot \epsilon$
- 6: $x_{t_{n-k}}^{\Psi, \omega} \leftarrow x_{t_n} + (1 + \omega) \cdot \Psi(x_{t_n}, t_n, t_{n-k}, c_1, c_2) - \omega \cdot \Psi(x_{t_n}, t_n, t_{n-k}, \emptyset, \emptyset)$
- 7: $\mathcal{L}(\theta, \theta^-; \Psi) \leftarrow d(f_\theta(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(\hat{x}_{t_{n-k}}, t_{n-k}, c_1, c_2, \omega))$
- 8: $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}(\theta, \theta^-)$
- 9: $\theta^- \leftarrow \text{stopgrad}(\mu \theta^- + (1 - \mu) \theta)$
- 10: **until** convergence

Compositional Consistency Distillation Algo

- Switch



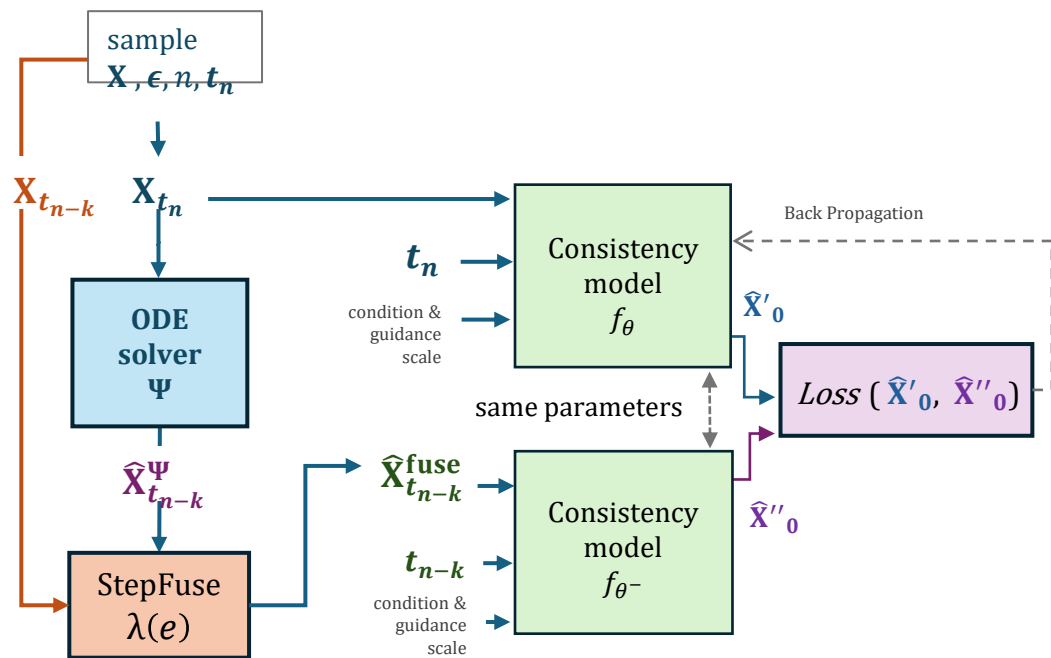
Algorithm 4.2 Modified Consistency Distillation–Switch

Input: dataset \mathcal{D} , initial model parameter θ , learning rate η , ODE solver $\Psi(\cdot, \cdot, \cdot, \cdot)$, distance metric $d(\cdot, \cdot)$, guidance scale $[\omega_{\min}, \omega_{\max}]$, skipping steps k , **switching threshold** e_{switch}

- 1: $\theta^- \leftarrow \theta$
- 2: **repeat**
- 3: Sample $(x, c_1, c_2) \sim \mathcal{D}$, $\omega \sim [\omega_{\min}, \omega_{\max}]$
- 4: Sample $\epsilon \sim \mathcal{N}(0, I)$, $n \sim \mathcal{U}[1 + k, N]$
- 5: $x_{t_n} \leftarrow \sqrt{\bar{\alpha}_{t_n}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_n}} \cdot \epsilon$
- 6: **if** $e < e_{\text{switch}}$ **then**
- 7: $x_{t_{n-k}}^{\Psi, \omega} \leftarrow x_{t_n} + (1 + \omega) \cdot \Psi(x_{t_n}, t_n, t_{n-k}, c_1, c_2) - \omega \cdot \Psi(x_{t_n}, t_n, t_{n-k}, \emptyset, \emptyset)$
- 8: $\mathcal{L}_{\text{Switch}}(\theta, \theta^-; \Psi) \leftarrow d\left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(x_{t_{n-k}}^{\Psi, \omega}, t_{n-k}, c_1, c_2, \omega)\right)$
- 9: **else**
- 10: $x_{t_{n-k}} \leftarrow \sqrt{\bar{\alpha}_{t_{n-k}}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_{n-k}}} \cdot \epsilon$
- 11: $\mathcal{L}_{\text{Switch}}(\theta, \theta^-) \leftarrow d\left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(x_{t_{n-k}}, t_{n-k}, c_1, c_2, \omega)\right)$
- 12: **end if**
- 13: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}_{\text{Switch}}(\theta, \theta^-)$
- 14: $\theta^- \leftarrow \text{stopgrad}(\mu \theta^- + (1 - \mu) \theta)$
- 15: **until** convergence

Compositional Consistency Distillation Algo

• StepFuse



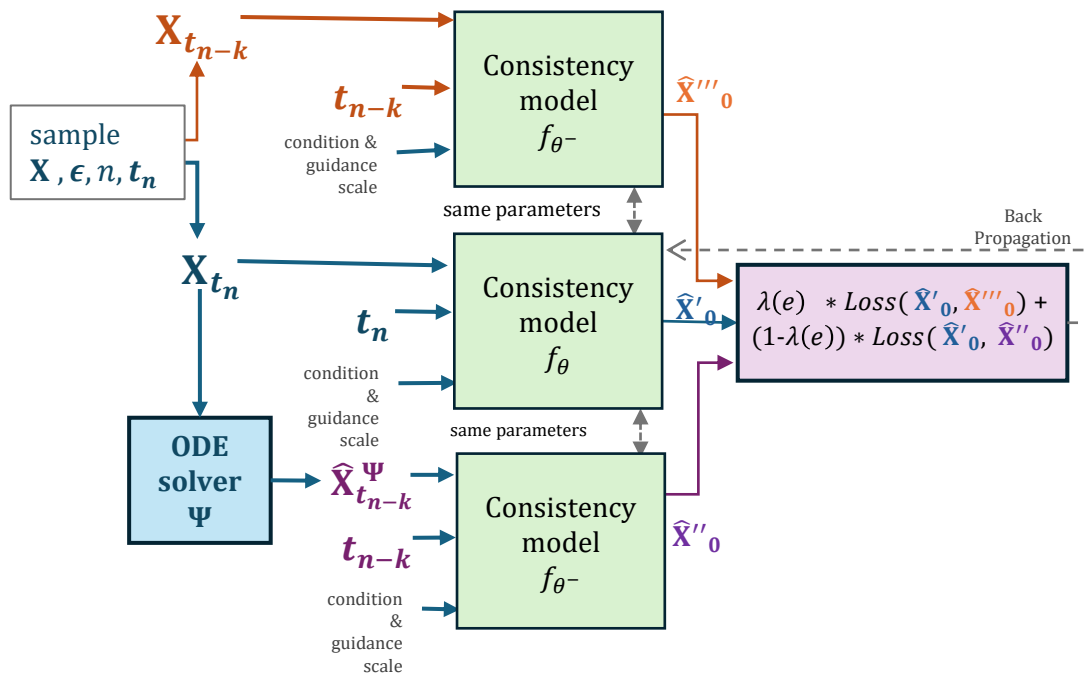
Algorithm 4.3 Modified Compositional label Consistency Distillation–StepFuse

Input: dataset \mathcal{D} , initial model parameter θ , learning rate η , ODE solver $\Psi(\cdot, \cdot, \cdot, \cdot)$, distance metric $d(\cdot, \cdot)$, guidance scale $[\omega_{\min}, \omega_{\max}]$, skipping steps k , **fuse scheduler** $\lambda(e)$

- 1: $\theta^- \leftarrow \theta$
- 2: **repeat**
- 3: Sample $(x, c_1, c_2) \sim \mathcal{D}$, and $\omega \sim [\omega_{\min}, \omega_{\max}]$
- 4: Sample $\epsilon \sim \mathcal{N}(0, I)$, $n \sim \mathcal{U}[1 + k, N]$
- 5: $x_{t_n} \leftarrow \sqrt{\bar{\alpha}_{t_n}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_n}} \cdot \epsilon$
- 6: $x_{t_{n-k}}^{\Psi, \omega} \leftarrow x_{t_n} + (1 + \omega) \cdot \Psi(x_{t_n}, t_n, t_{n-k}, c_1, c_2) - \omega \cdot \Psi(x_{t_n}, t_n, t_{n-k}, \emptyset, \emptyset)$
- 7: $x_{t_{n-k}} \leftarrow \sqrt{\bar{\alpha}_{t_{n-k}}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_{n-k}}} \cdot \epsilon$
- 8: $\hat{x}_{t_{n-k}}^{\text{fuse}} \leftarrow \lambda(e) \cdot \hat{x}_{t_{n-k}}^{\Psi, \omega} + (1 - \lambda(e)) \cdot x_{t_{n-k}}$
- 9: $\mathcal{L}_{\text{Stepfuse}}(\theta, \theta^-; \Psi) \leftarrow d(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(\hat{x}_{t_{n-k}}^{\text{fuse}}, t_{n-k}, c_1, c_2, \omega))$
- 10: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}_{\text{Stepfuse}}(\theta, \theta^-)$
- 11: $\theta^- \leftarrow \text{stopgrad}(\mu \theta^- + (1 - \mu) \theta)$
- 12: **until** convergence

Compositional Consistency Distillation Algo

• LossFuse



Algorithm 4.4 Modified Compositional label Consistency Distillation–LossFuse

Input: dataset \mathcal{D} , initial model parameter θ , learning rate η , ODE solver $\Psi(\cdot, \cdot, \cdot, \cdot)$, distance metric $d(\cdot, \cdot)$, guidance scale $[\omega_{min}, \omega_{max}]$, skipping steps k , **fuse scheduler** $\lambda(e)$

- 1: $\theta^- \leftarrow \theta$
- 2: **repeat**
- 3: Sample $(x, c_1, c_2) \sim \mathcal{D}$, and $\omega \sim [\omega_{min}, \omega_{max}]$
- 4: Sample $\epsilon \sim \mathcal{N}(0, I)$, $n \sim \mathcal{U}[1 + k, N]$
- 5: $x_{t_n} \leftarrow \sqrt{\bar{\alpha}_{t_n}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_n}} \cdot \epsilon$
- 6: $x_{t_{n-k}}^{\Psi, \omega} \leftarrow x_{t_n} + (1 + \omega) \cdot \Psi(x_{t_n}, t_n, t_{n-k}, c_1, c_2) - \omega \cdot \Psi(x_{t_n}, t_n, t_{n-k}, \emptyset, \emptyset)$
- 7: $x_{t_{n-k}} \leftarrow \sqrt{\bar{\alpha}_{t_{n-k}}} \cdot x + \sqrt{1 - \bar{\alpha}_{t_{n-k}}} \cdot \epsilon$
- 8: $\mathcal{L}_{teacher}(\theta, \theta^-; \Psi) \leftarrow d\left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(\hat{x}_{t_{n-k}}^{\Psi, \omega}, t_{n-k}, c_1, c_2, \omega)\right)$
- 9: $\mathcal{L}_{fpf}(\theta, \theta^-) \leftarrow d\left(f_{\theta}(x_{t_n}, t_n, c_1, c_2, \omega), f_{\theta^-}(x_{t_{n-k}}, t_{n-k}, c_1, c_2, \omega)\right)$
- 10: $\mathcal{L}_{lossfuse} \leftarrow \lambda(e) \cdot \mathcal{L}_{teacher} + (1 - \lambda(e)) \cdot \mathcal{L}_{fpf}$
- 11: $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}_{lossfuse}(\theta, \theta^-)$
- 12: $\theta^- \leftarrow \text{stopgrad}(\mu \theta^- + (1 - \mu) \theta)$
- 13: **until** convergence

Experiment setup

- Training Configuration

Hyperparameters fixed across all experiments

Training timesteps	1000
Batch Size	24
Noise scheduler β	0.0001-0.2 linear
Base channels dimension	128
# of Residual Blocks	2
Channel Multiplier	1,2,4,4
Optimizer	AdamW
ω embedding dimension	128
Huber loss param	$\delta = 0.001$
EMA decay	$\mu = 0.995$